

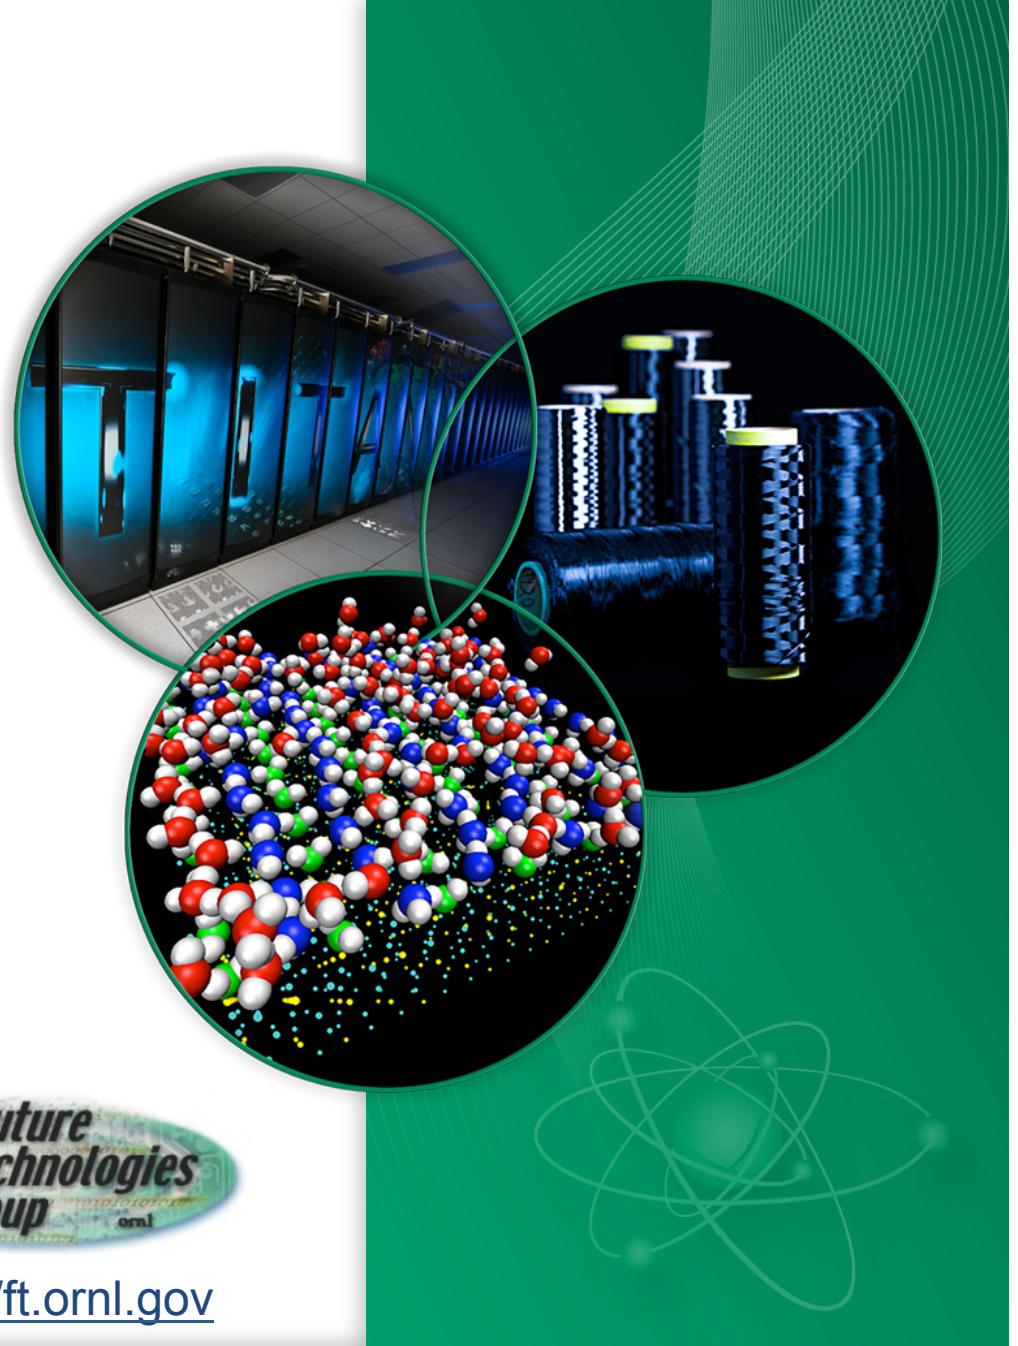
OpenARC: Open Accelerator Research Compiler (mini-workshop)

Seyong Lee

Future Technologies Group

Oak Ridge National Laboratory

lees2@ornl.gov



<http://ft.ornl.gov>

Agenda

- 9:00 – 10:20 am Introduction to OpenARC
- 10:20 – 10:50 am How to use OpenARC
(User Guide)
- 10:50 – 11:00 am Breaktime
- 11:00 – noon OpenARC Internals
(Advanced Topics)
- 1:00 – 2:00 pm Introduction to OpenACC
- 2:00 – 4:00 pm OpenACC and OpenARC
Experience

Motivation

- Scalable Heterogeneous Computing (SHC)
 - Enabled by graphics processors, Intel Xeon Phi, or other non-traditional devices.
 - Emerging solution to respond to the constraints of energy, density, and device technology trends.
 - However, the complexity in SHC systems causes portability and productivity issues.

Directive-Based Accelerator Programming Models

- Provide abstraction over architectural details and low-level programming complexities.
- However, too much abstraction puts significant burdens on:
 - Performance tuning
 - Debugging
 - Scaling
- We need in-depth evaluation and research on the directive-based, heterogeneous programming to address the two conflicting goals in SHC systems: productivity and portability.

Related Work

- Custom-Directive-based GPU Compiler
 - hiCUDA, OpenMPC
 - PGI accelerator compiler, HMPP, Rstream, etc.
- OpenACC compiler
 - PGI compiler, CAPS compiler, Cray compiler, etc.
 - accULL, OpenUH
- OpenMP4 compiler for accelerator computing
 - HOMP, GCC

OpenARC: Open Accelerator Research Compiler

- Open-Sourced, High-Level Intermediate Representation (HIR)-Based, Extensible Compiler Framework.
 - Perform source-to-source translation from OpenACC C to target accelerator models.
 - Support full features of OpenACC V1.0 (+ subset of OpenACC V2.0 + array reductions and function calls)
 - Support both CUDA and OpenCL as target accelerator models (NVIDIA GPUs, AMD GPUs, Xeon Phi, etc.)
 - OpenMP V4.0 support is being tested.
 - Provide common runtime APIs for various back-ends
 - Can be used as a research framework for various study on directive-based accelerator computing.
 - Built on top of Cetus compiler framework, equipped with various advanced analysis/transformation passes and built-in tuning tools.
 - OpenARC's IR provides an AST-like syntactic view of the source program, easy to understand, access, and transform the input program.

Design Goals of OpenARC

- Extensibility
 - Extensible OpenARC IR class hierarchy:
 - Its IR class hierarchy is easily extended to embrace new language constructs.
 - Rich semantic annotations:
 - Each OpenARC IR object can be augmented with rich annotations.
 - User can create different annotation types by deriving from a base PragmaAnnotation class.
 - Common, abstract OpenARC runtime API:
 - Unified interface between the program and the target-specific runtime

Overall Design of OpenARC (cont.)

- Debuggability
 - AST-like HIR, combined with abstract OpenARC runtime, allows to generate output codes similar to input program.
 - Clear separation between analyses and transformations
 - All analysis outputs are stored as internal annotations, used as input to transformation passes → Easy to build traceability mechanisms

Design Goals of OpenARC (cont.)

- Tunability
 - Provide a rich set of directives to control various compiler optimizations and hardware-specific features.
 - Its built-in tuning tools combined with the rich OpenARC directives allow users to control overall OpenACC-to-Accelerator translation in a fine-grained, but still abstract manner.

OpenARC Directive Extension and Environment Variables

- OpenARC Directives

#pragma omp openmp-directive [clause [,] clause]...]

#pragma acc openacc-directive [clause [,] clause]...]

#pragma cetus [clause[,] clause]...]

#pragma acc internal [clause [,] clause]...]

#pragma openarc cuda [clause [,] clause]...]

#pragma openarc resilience-directive [clause [,] clause]...]

#pragma openarc profile-directive [clause [,] clause]...]

#pragma aspen aspen-directive [clause [,] clause]...]

- OpenARC Environment Variables

- Control the program-level behavior of various optimizations or execution configurations for an output GPU program.

OpenARC Runtime API

- HeteroIR: High-Level, Architecture-Independent Intermediate Representation
 - Used as an intermediate language to map high-level programming models (OpenACC) to diverse heterogeneous devices.
- Primary Constructs in HeteroIR (Partial List)

 Configuration

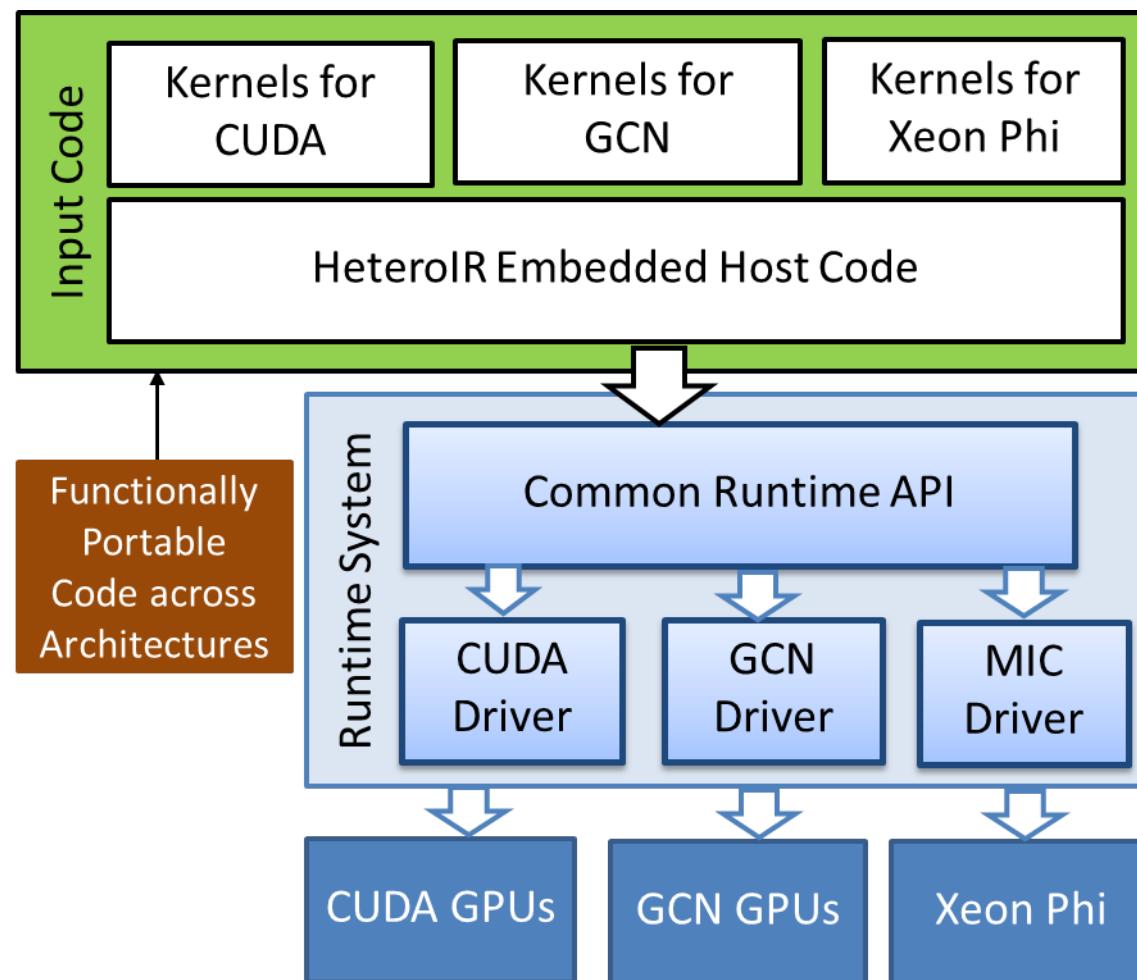
 Kernels

 Memory

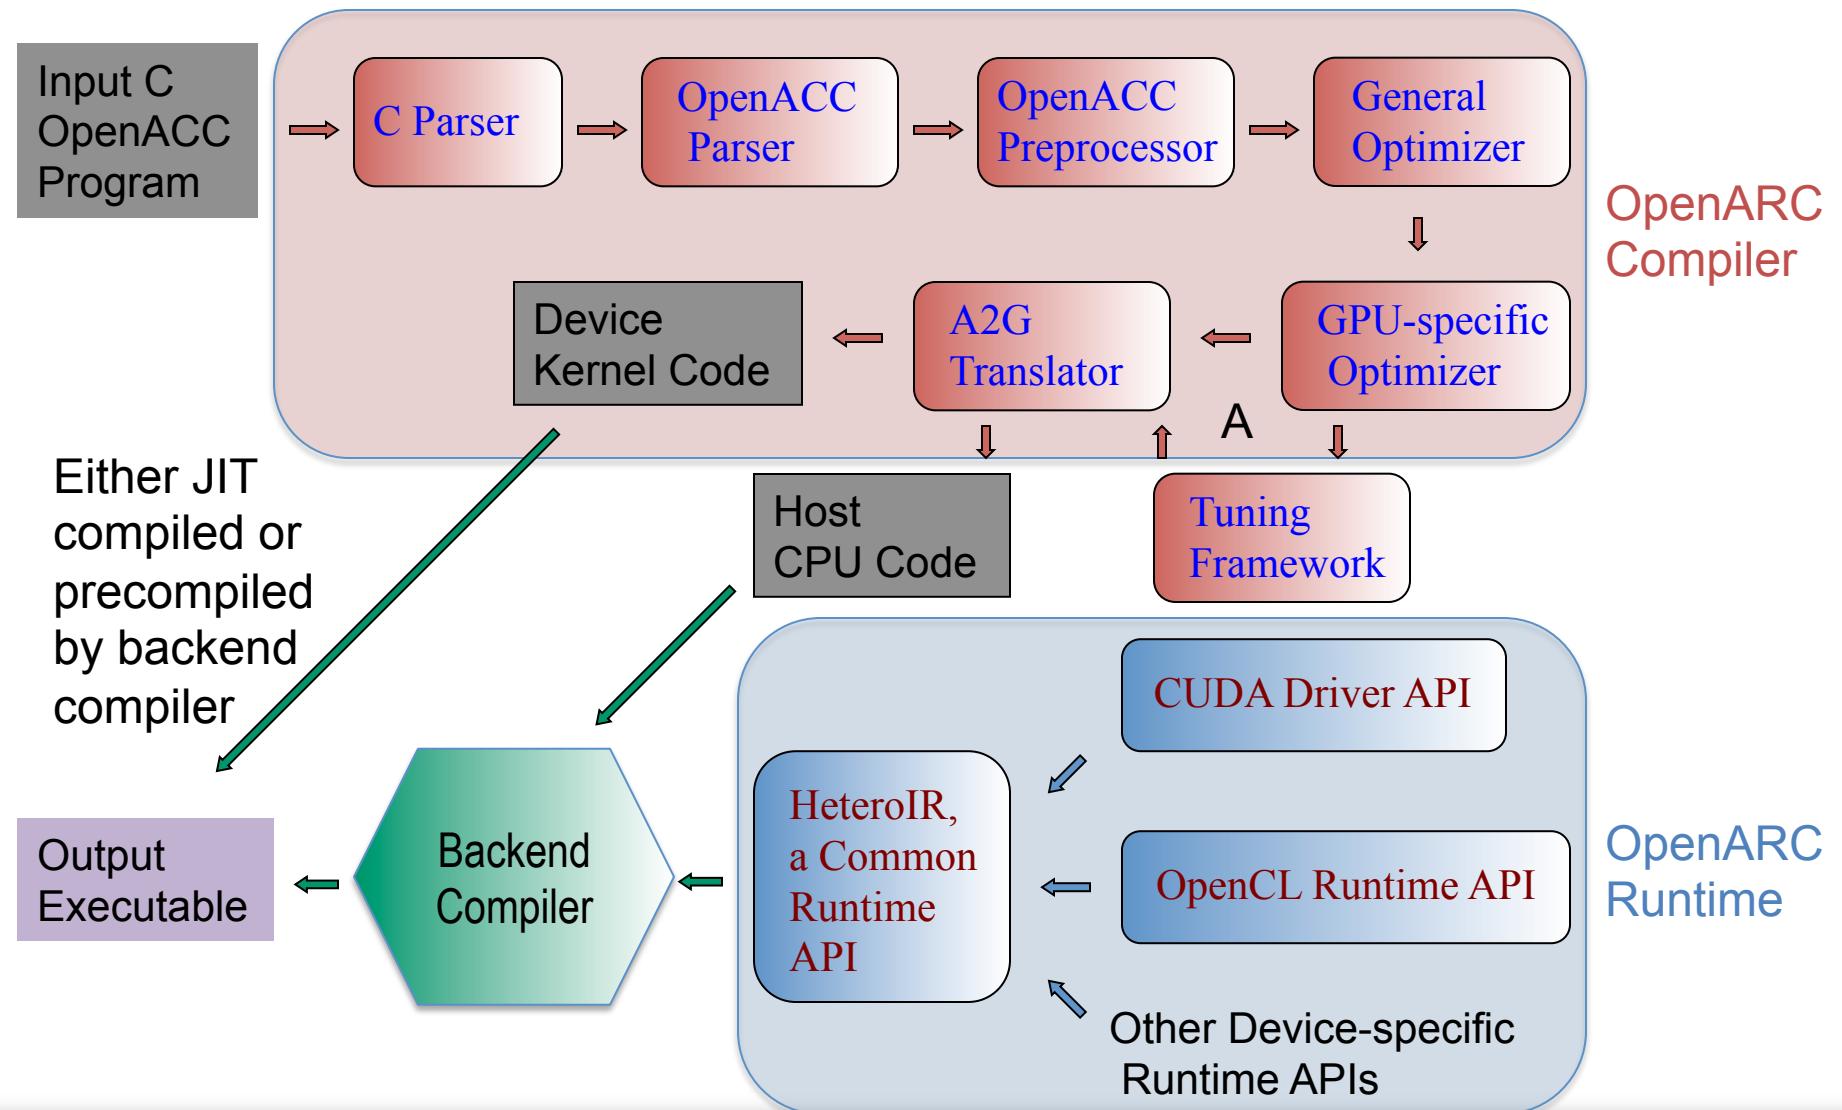
 Synchronization

HI_set_device (device type, device number)	Set the device to use
HI_init ()	Initialize the previously selected device.
HI_reset ()	Deinitialize the device.
HI_malloc (pointer, size, queue)	Allocate memory on the device.
HI_memcpy (source, destination, size, queue)	Perform synchronous/asynchronous data transfers.
HI_get_device_address (device address, host address)	Check if the host address has a valid mapping on the device. Return the device address if mapping is present.
HI_free (pointer, queue)	Free the device memory pointed to by the pointer.
HI_register_kernel_arg (kernel name, parameter, ...)	Attach the argument to the corresponding kernel.
HI_kernel_call (kernel name, grid size, queue)	Launch the kernel with the specified grid size on the specified queue.
HI_set_async (queue)	Set the specified queue number.
HI_async_wait (queue)	Wait until all actions on the specified queue are finished.
HI_async_wait_all ()	Wait for all queues to finish the actions queued on them.

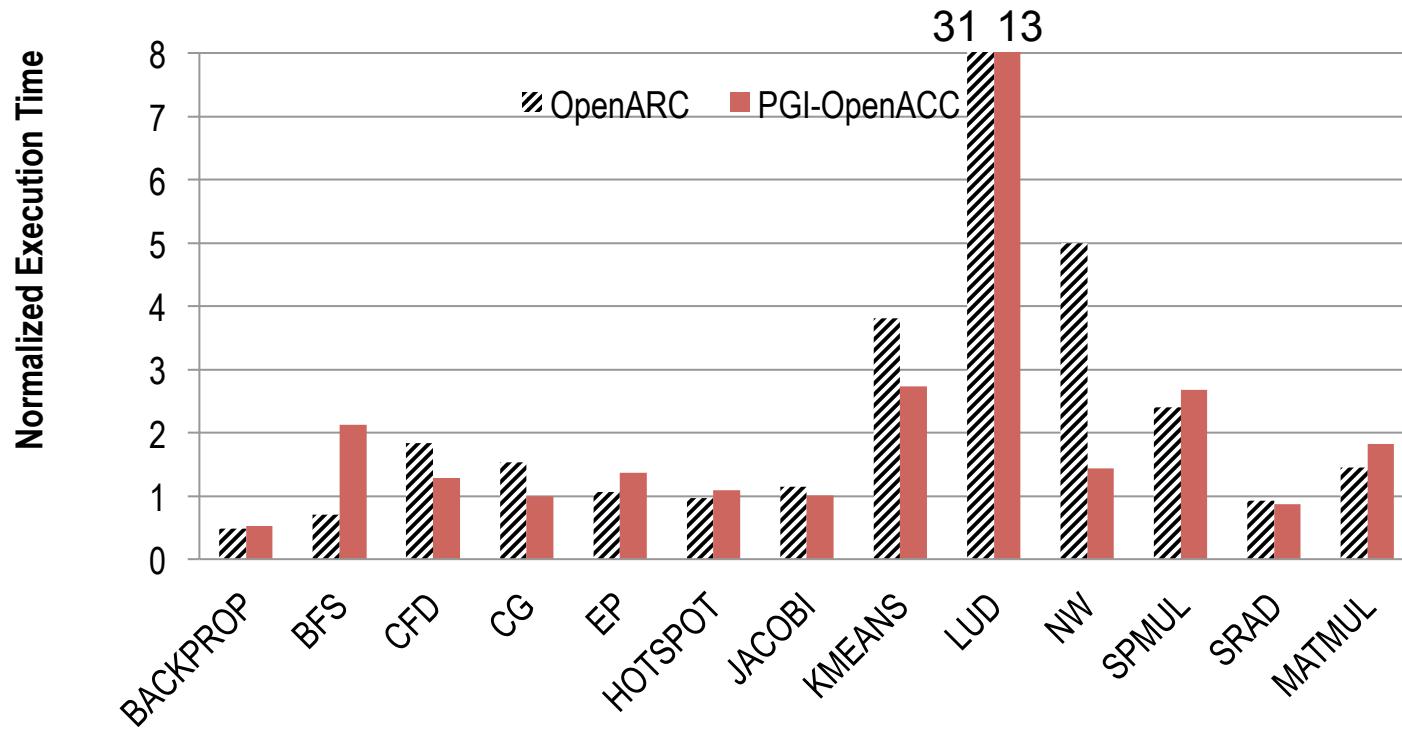
HeteroIR-based, OpenARC Runtime System



OpenARC System Architecture

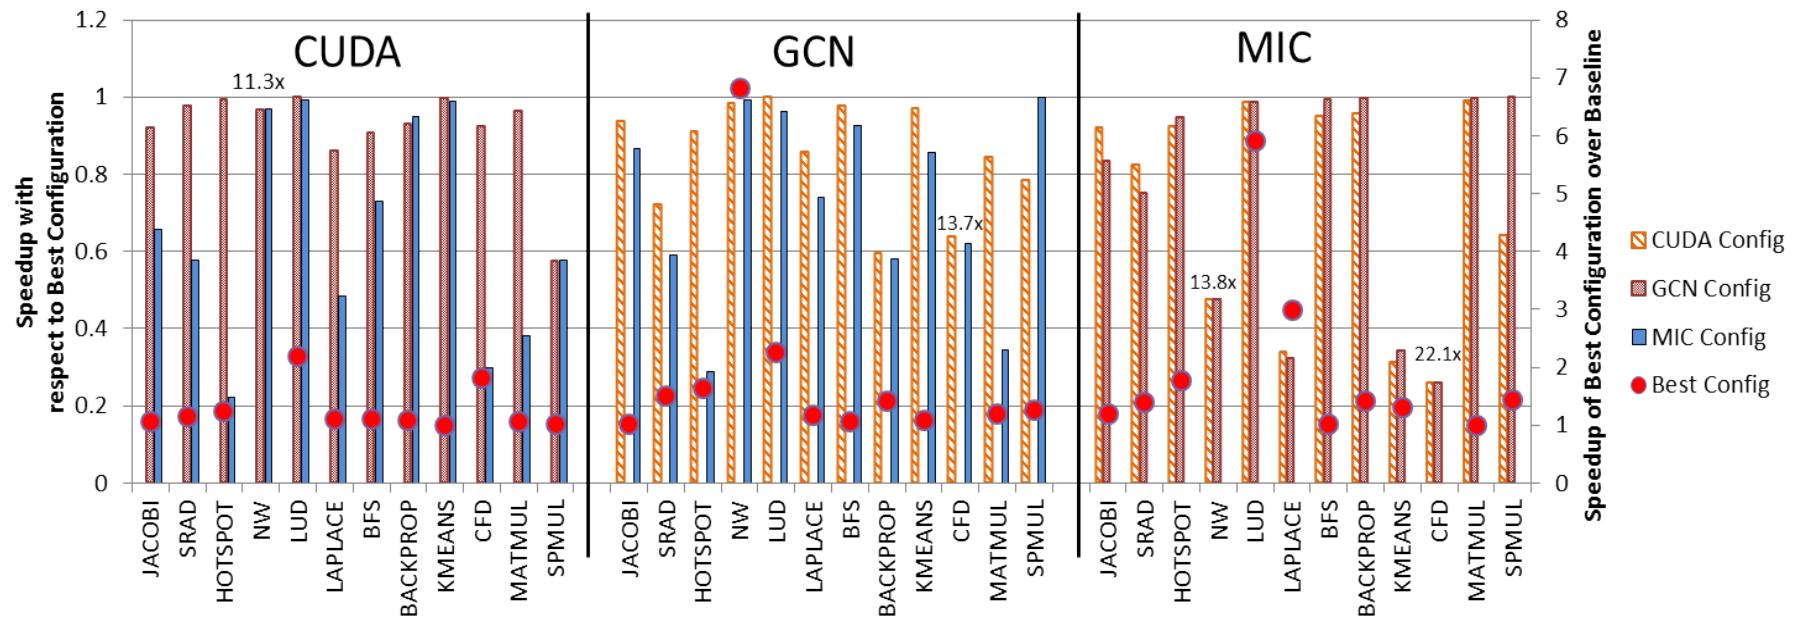


Preliminary Evaluation of OpenARC



Performance of OpenARC and PGI-OpenACC compiler (V13.6)
relative to manual CUDA versions (Lower is better.)

Preliminary Evaluation of OpenARC (2)



Performance Portability Evaluation: Best performing OpenACC on one architecture may not produce best performance on another architecture. Tuning can be highly beneficial over the baseline translated code.

Research Agenda

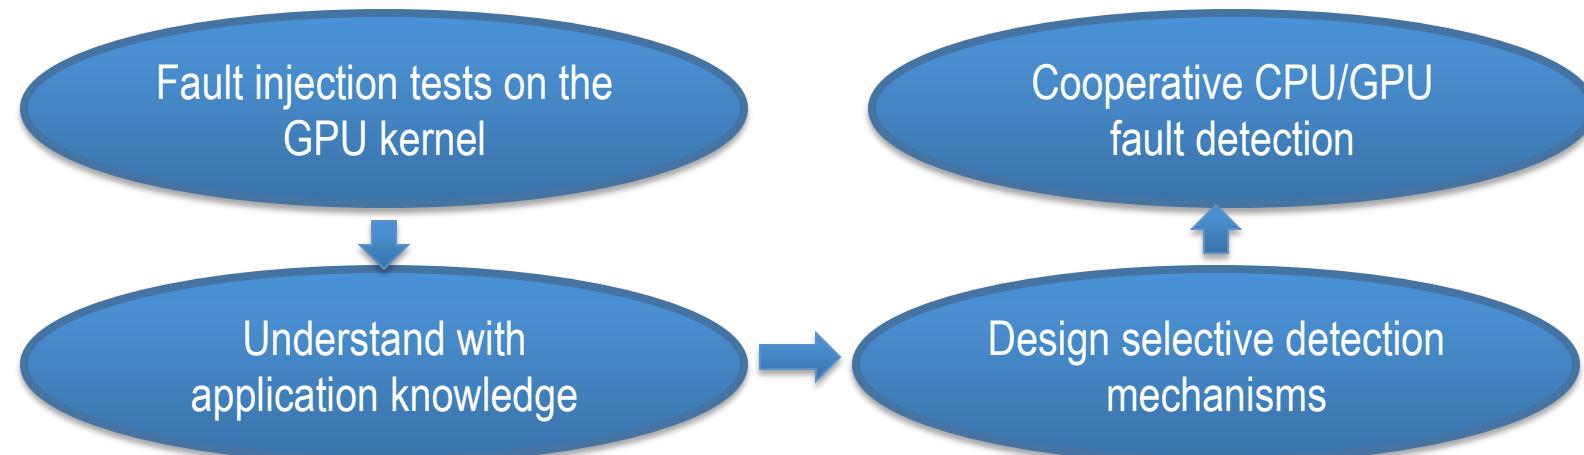
- Research In Progress
 - Directive-based, interactive program debugging and optimization
 - Directive-based, application-driven cooperative CPU/GPU fault detection
 - autoASPEN: directive-based, automatic performance modeling
 - Directive-based, scalable heterogeneous computing
 - Compiler-driven, custom profiling system for directive programming
 - OpenACC-e: OpenACC extension to support accelerator-specific features at high-level

Research in Progress: Directive-Based, Interactive Program Debugging and Optimization

- Problem
 - *Too much abstraction* in directive-based GPU programming!
 - Debuggability: difficult to diagnose logic errors and performance problems at the directive level
 - Performance Optimization: difficult to find where and how to optimize
 - Solution
 - Directive-based, interactive GPU program debugging and optimization
 - OpenARC compiler:
 - Generates runtime codes necessary for *GPU-kernel verification* and *memory-transfer verification and optimization*.
 - Runtime
 - Locate trouble-making kernels by comparing execution results at kernel granularity.
 - Trace the runtime status of CPU-GPU coherence to detect incorrect/missing/redundant memory transfers.
 - Users
 - Iteratively fix/optimize incorrect kernels/memory transfers based on the runtime feedback and apply to input program.
-
- The diagram illustrates the iterative process of the OpenARC system. It features three main components: 'User' (blue rounded rectangle), 'Runtime' (blue rounded rectangle), and 'OpenARC' (blue rounded rectangle). A circular arrow connects the 'User' and 'Runtime' components. Another circular arrow connects the 'Runtime' and 'OpenARC' components. The 'OpenARC' component is positioned above the 'Runtime' component, indicating its role as a compiler or interface between the user and the runtime environment.
- Iteratively find where and how to fix/optimize**

Research in Progress (2): Directive-Based, Application-Driven Cooperative CPU/GPU Fault Detection

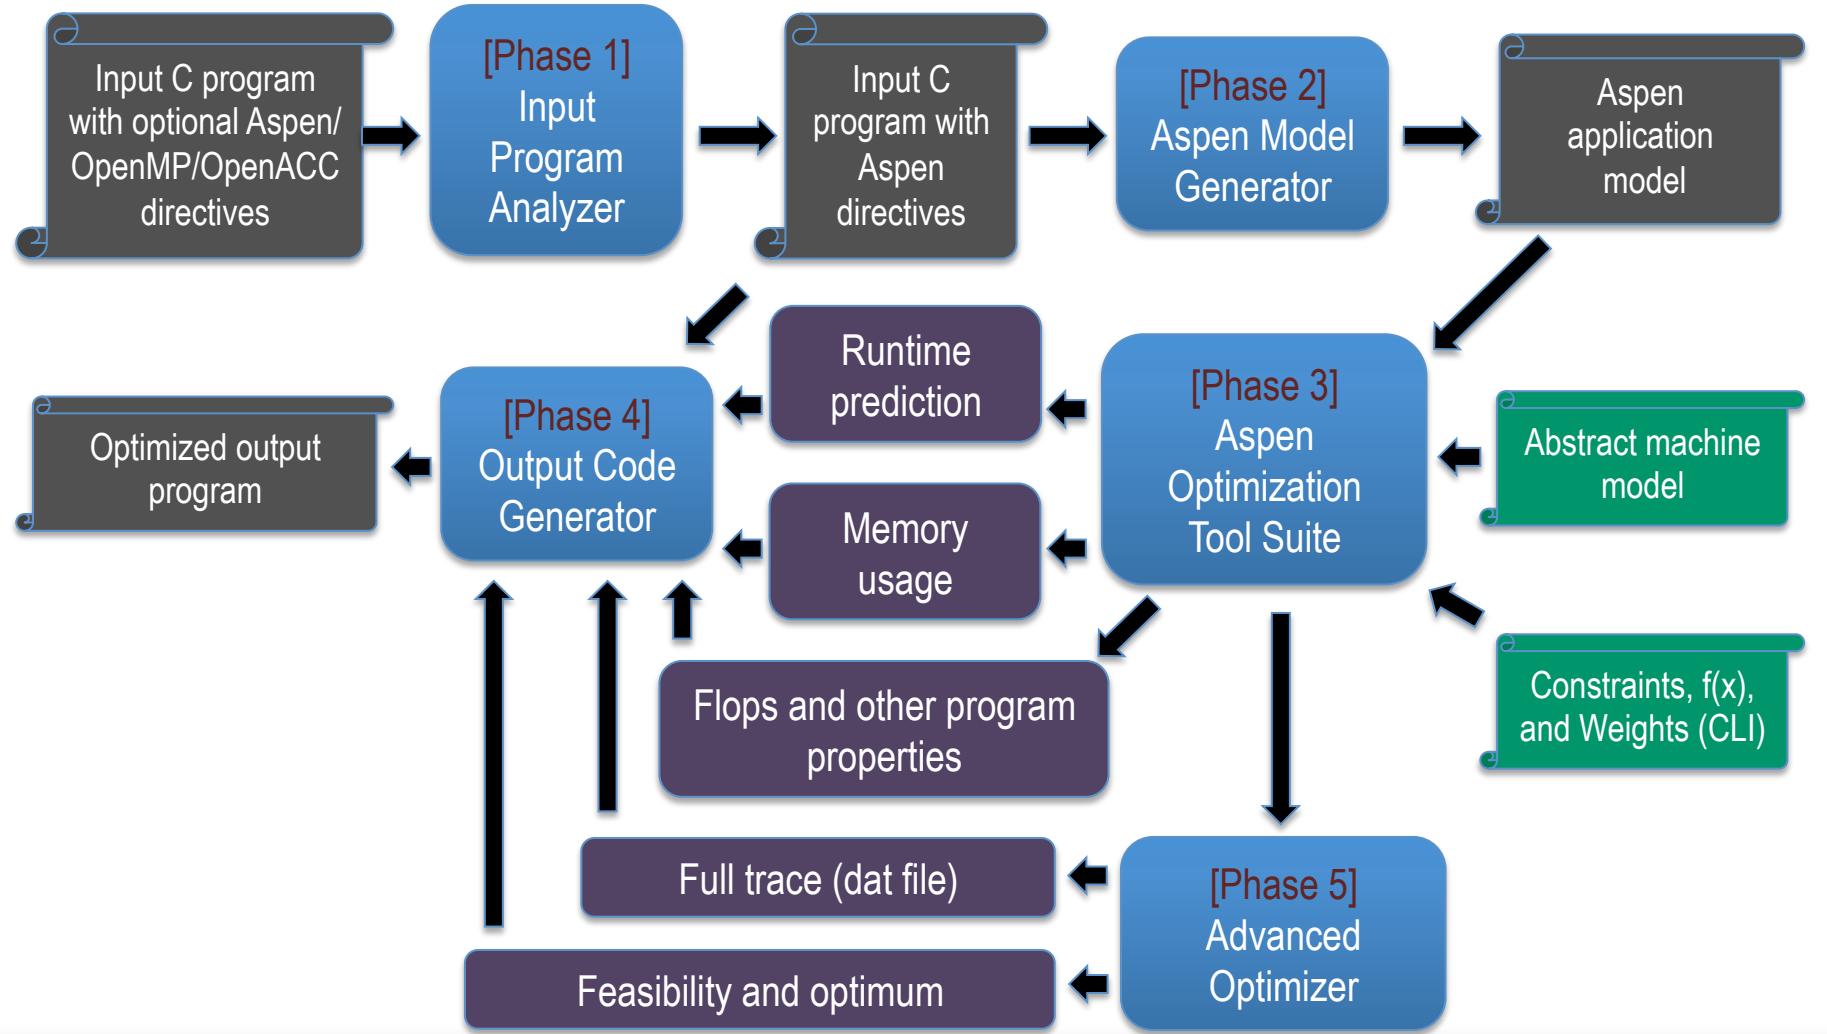
- Problem
 - *Limitations of the conventional resilience mechanisms*
 - No cooperation for data protection between CPU and GPU
 - Enforce a uniform protection, regardless of application semantics.
- Solution
 - Explore the viability of using an application-driven CPU/GPU cooperative method.
 - A user-level fault injection framework, which uses a directive-based approach to configure fault injection
 - Application-driven, cooperative CPU/GPU fault detection



Research in Progress (3): autoAspen: Directive-Based, Automatic Performance Modeling

- Problem
 - Porting programming systems (compilers, runtime, and libraries) across architectures is increasingly critical as architectures grow more complex and change quickly.
 - Many of performance parameters for the target architecture are implicitly embedded in the programming systems; hard to move a programming system for one configuration to another.
- Solution
 - Propose a method and system for performance portability in programming systems that separates performance parameters of each architecture from the implementation of the programming system.
 - Generate an explicit performance model from the application source code.
 - The parameters for the target architecture are convolved with the performance model to allow intelligent and portable decision by the programming system.

Research in Progress (3): autoAspen: Process Diagram of The Overall Design Exploration



Summary

- OpenARC is an open-sourced, High-level Intermediate Representation (HIR)-based, extensible compiler framework.
- HIR with a rich set of directives in OpenARC provides a powerful research framework for various source-to-source translation and instrumentation experiments.
- The additional OpenARC directives with its built-in tuning tools allow users to control overall OpenACC-to-GPU translation in a fine-grained, but still abstract manner.

Session 2: How to Use OpenARC (User Guide)

- Getting Started
 - Requirements
 - Installation
 - Environment setup and running OpenARC
 - Known Issues
- OpenARC Applications
 - Performance Tuning in OpenARC
 - Directive-based, interactive program debugging and optimization
 - Directive-based, application-driven fault injection
 - autoASPEN: directive-based, automatic performance modeling

Getting Started

- Requirements
 - JAVA SE 6 or later
 - GCC
 - ANTLRv2
 - Default antlr.jar file is included in the OpenARC distribution.
- Installation
 - Obtain OpenARC distribution
 - No public distribution is available yet.
 - Internal Git repository: <https://code.ornl.gov/f6l/OpenARC.git>
 - Contact us to access the source code (lees2@ornl.gov)

Getting Started (2)

- Installation
 - Build OpenARC compiler:
 - Use provided “build.sh” or “build.xml” files.
\$ build.sh jar //create jar file of OpenARC.
 - Build OpenARC runtime:
 - Go to ./openarcrt directory
 - Copy “make.header.sample” to “make.header”.
 - Update “make.header” as necessary.
 - Run “batchmake.bash” script
\$ batchmake.bash

Running OpenARC

- Environment Setup
 - Set OPENARC_ARCH to 0 for CUDA (default), 1 for OpenCL, or 2 for OpenCL for Xeon Phi.
 - Set OpenACC environment variables (e.g., ACC_DEVICE_TYPE, ACC_DEVICE_NUM, etc.) to use non-default configuration.
 - Set OMP_NUM_THREADS to the maximum number of OpenMP threads if OpenMP is used in the input OpenACC program.
 - Set OPENARC_JITOOPTION to pass options to the backend runtime compiler (e.g., NVCC options for JIT CUDA kernel compilation or clBuildProgram options for JIT OpenCL kernel compilation.)
 - Set OPENARCRT_VERBOSITY to set the verbosity level (>= 0) of the OpenARC runtime in the profile mode.
- Running OpenARC

```
$ java -classpath=<user_class_path> openarc.exec.ACC2GPUDriver  
<OpenARC-options> <Input C files>
```

OpenARC commandline options

- To find a list of available OpenARC options, use “-help” or “-dump-options”
\$ java openacc.exec.ACC2GPUDriver –dump-options
- To pass multiple options in one file, use “-gpuConfFile=confFile” option.
- Some Important options
 - macro: used to pass macros to the C preprocessor.
 - addIncludePath: used to pass paths for header files.
 - ... (many more optional ones)

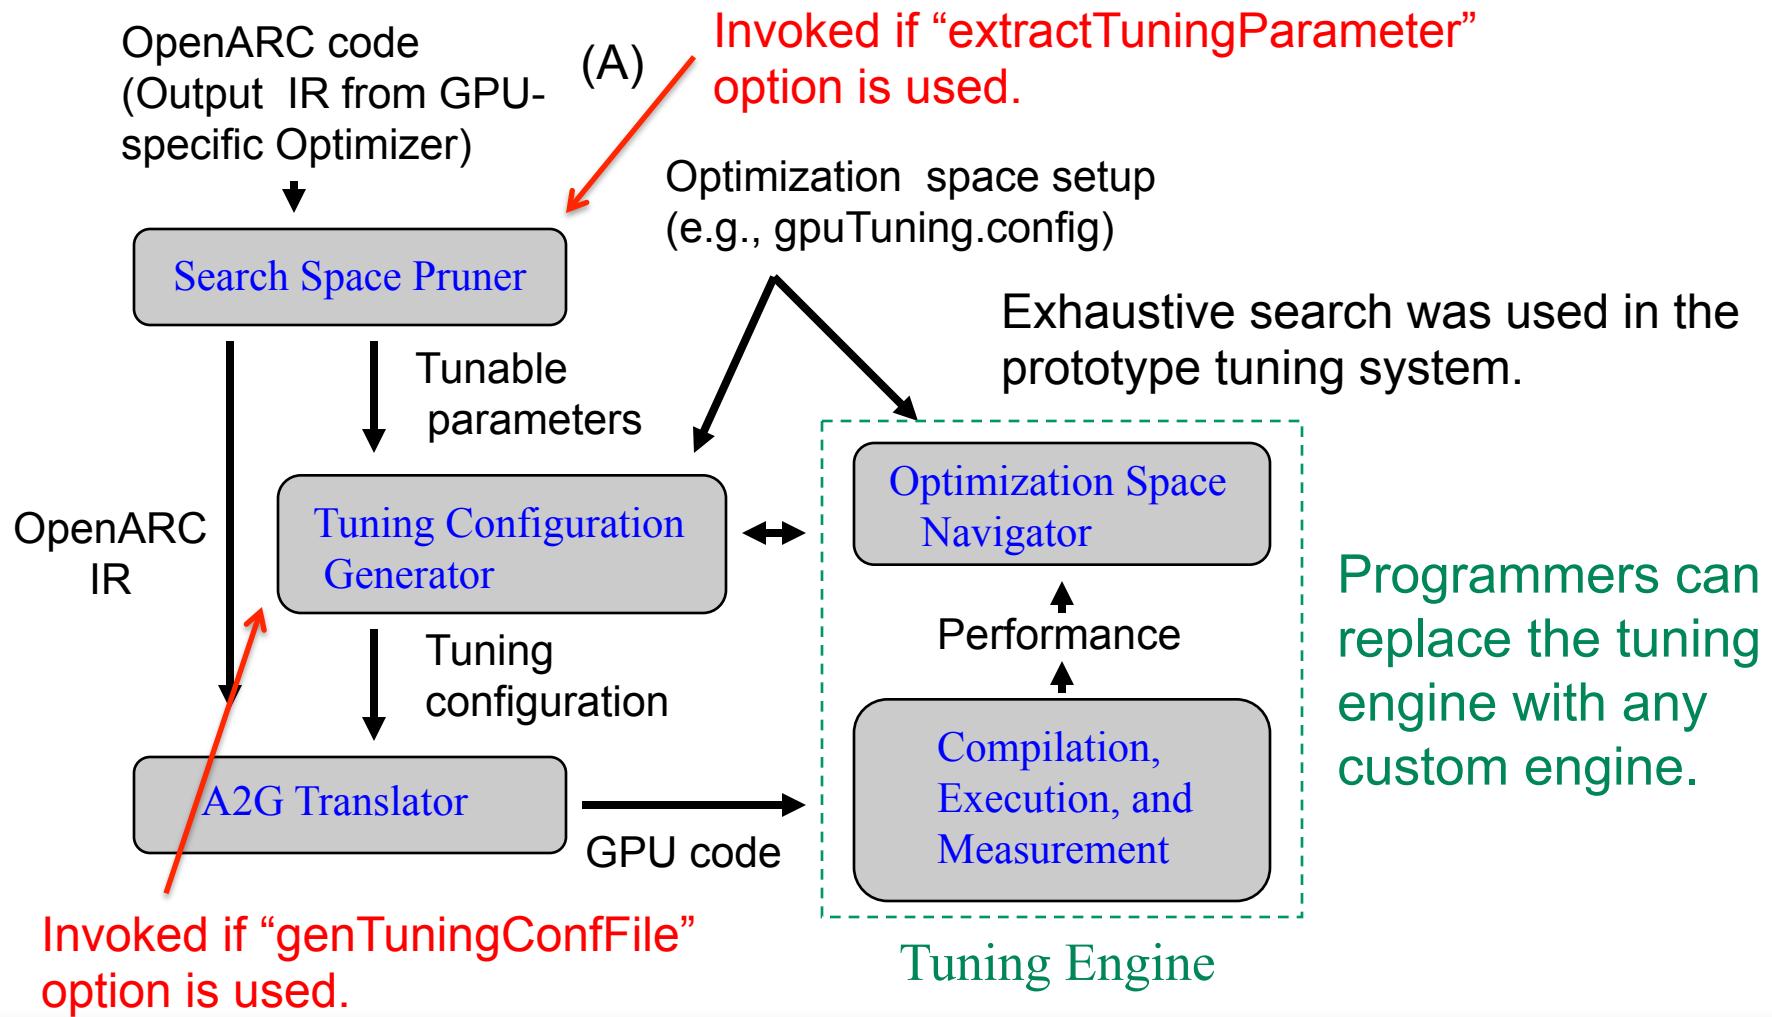
Known Issues on the Current OpenARC Implementation

- The C parser supports C99 features only partially.
 - E.g., mixed declaration and code, C99 attributes, and C99 macros are only partially supported.
- GNU C preprocessor does not expand macros in pragma annotation.
 - To expand in annotations, use “#pragma acc #define ...” directive.
- Struct member is not allowed in OpenACC data clauses.
- Partial array passing is allowed only if the start index is 0.
- OpenACC vector clauses are ignored.
- OpenACC data regions can have compute regions in called functions, but a compute region can have a function call only if the called function does not contain worksharing loops.

Session 2: How to Use OpenARC (User Guide)

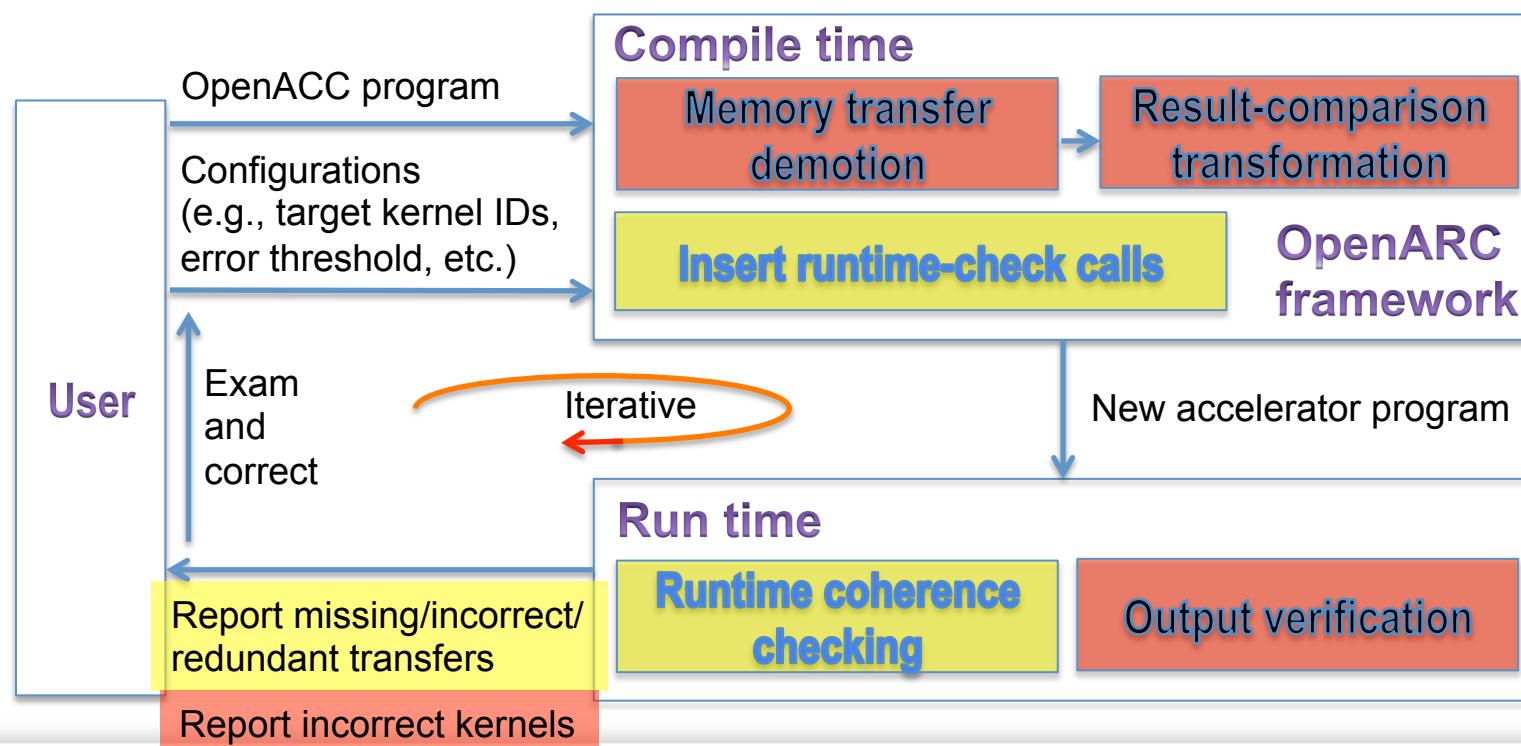
- Getting Started
 - Requirements
 - Installation
 - Environment setup and running OpenARC
 - Known Issues
- OpenARC Applications
 - Performance Tuning in OpenARC
 - Directive-based, interactive program debugging and optimization
 - Directive-based, application-driven fault injection
 - autoASPEN: directive-based, automatic performance modeling

OpenARC Application: Performance Tuning using a built-in Tuning Framework in OpenARC



OpenARC Application 2: Directive-Based Interactive Program Debugging and Optimization

- Kernel verification
(invoked if “programVerification=2” option is used.)
- Memory-transfer verification and optimization
(invoked if “programVerification=1”)



OpenARC Application 3: Directive-Based, Application-Driven Fault Injection

- We develop a user-level fault injection framework for general C programs.
 - Propose a set of user-directives and C library functions.
 - OpenARC automatically inserts fault injection codes according to the user-directives.
- We fully emulate the randomness of fault occurrences.
 - Temporal randomness
 - Spatial randomness
- We use a CPU/GPU cooperative approach for emulating random fault injection

Directive-Based, Application-Driven Fault Injection (2)

- The directives controls fault injection configuration
 - Which program variable at which target region?
 - The condition to enable faults?
 - Number of fault injections and number of faulty bits
- Directive Format (partial list)

```
#pragma acc resilience [clause[,] clause]...]
```

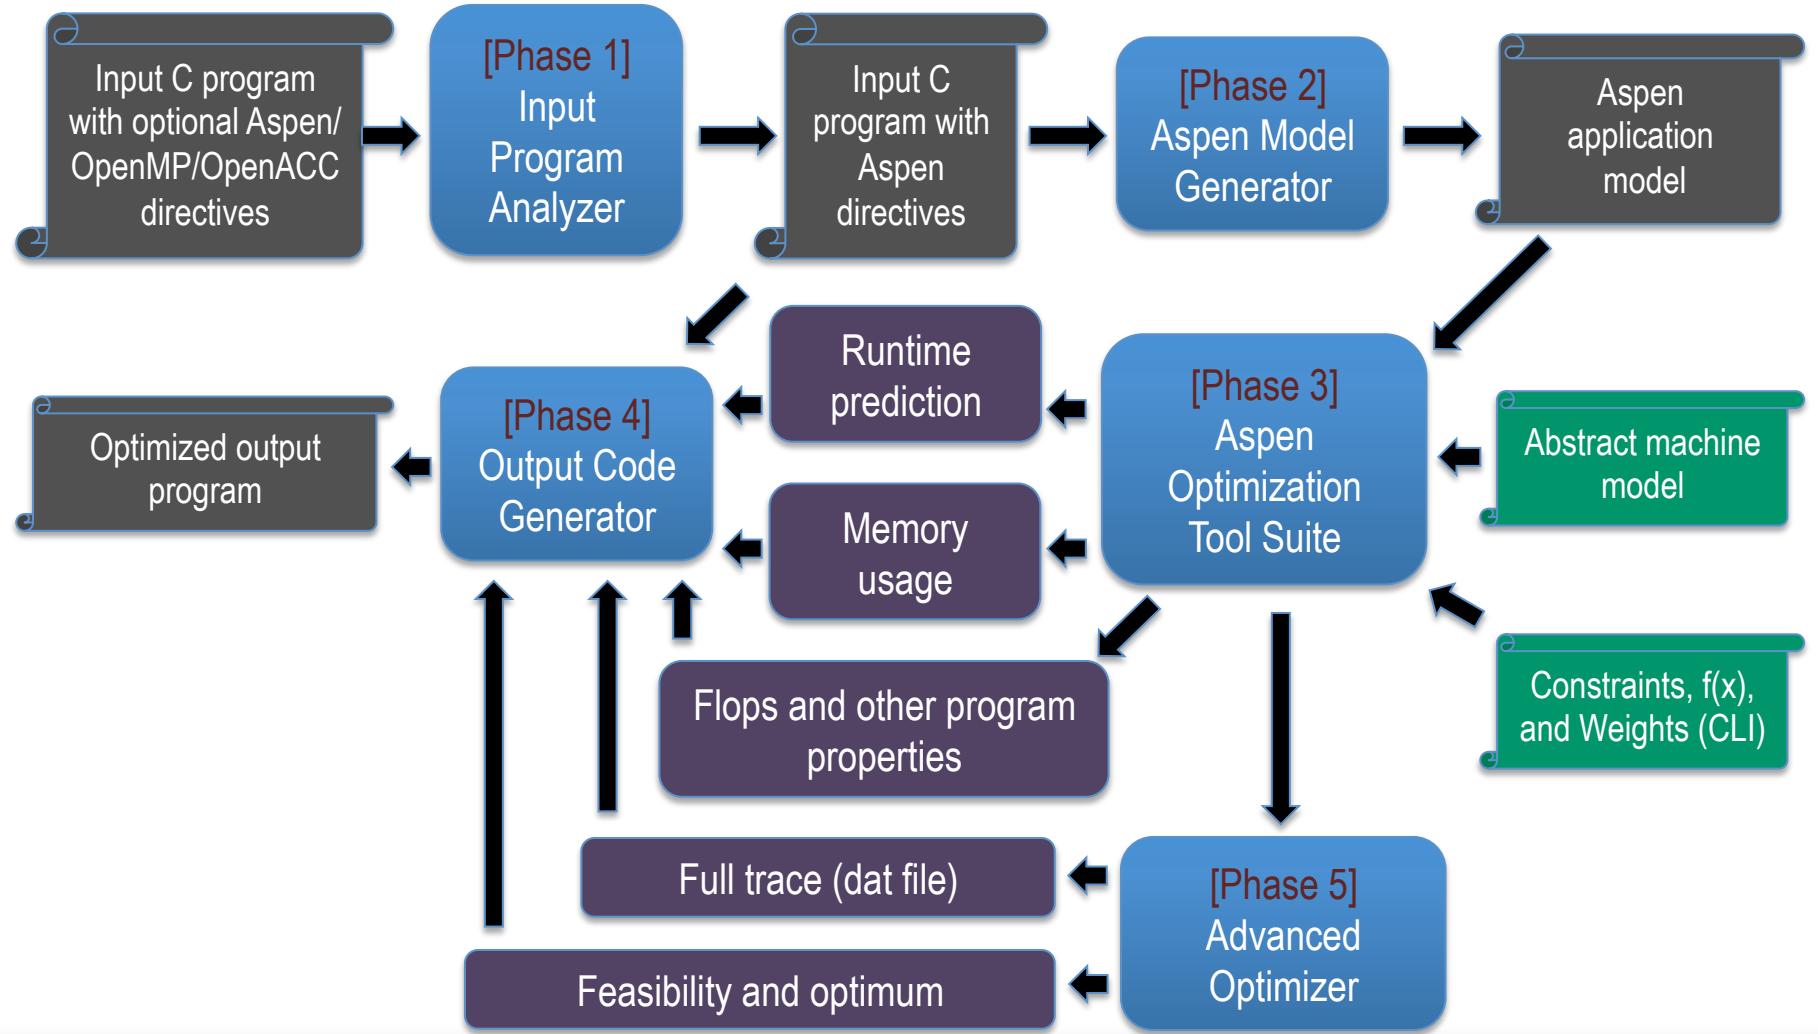
structured block

```
// where clause is one of the following: ftregion, ftcond(condition), ftdata(list),  
num_faults(exp), num_ftbits(exp), ftthread(exp)
```

```
#pragma acc ftinject [clause[,] clause]...]
```

```
// where clause is one of the following: ftdata(list), ftthread(exp)
```

OpenARC Application 4: autoAspen: Directive-Based, Automatic Performance Modeling



autoAspen: Directive-Based Automatic Performance Modeling (2)

- The proposed Aspen directives are primarily used by OpenARC to annotate important properties of the target program, such as computations, memory accesses, communications, control flows, etc.
- Aspen directives allow users to control how OpenARC analyzes and generates a performance model in an interactive manner.
- Aspen Modeling Directive Format
 - `#pragma aspen enter modelregion [label(name)]`
 - `#pragma aspen exit modelregion [label(name)]`
 - `#pragma aspen modelregion [label(name)]`
 - `#pragma aspen declare [clause [,] clause...]`
 - `#pragma aspen control [clause [,] clause...]`

Breaktime

- It's time for a break!
 - We will resume at 11:00AM

Session 3: OpenARC Internals (Advanced Topics)

- OpenARC Implementation
- OpenARC Internal Representation (IR)
 - OpenARC IR class hierarchy
 - Working on OpenARC IR
 - Example Transformations
- Built-in Cetus Passes for Program Analysis and Transformation

Implementation of OpenARC

- Built on top of the Cetus compiler infrastructure
- What is Cetus?
 - Source-to-source C compiler written in Java and maintained at Purdue University (<http://cetus.ecn.purdue.edu>)
 - Provides an internal C parser (Antlr, 23K+ lines)
 - Intermediate representations (IR) (23K+ lines)
 - Compiler passes (45K+ lines and growing)
 - Contains various compile-time analysis/transformation passes, including privatization, reduction recognition, symbolic expression manipulators, induction variable substitution, etc.
- OpenARC adds 65K+ lines for OpenACC-to-Accelerator translation.

OpenARC Source Structure

- OpenARC source

Cetus.base | C Parser

Cetus.hir Openacc.hir | OpenARC IR

Cetus.analysis Openacc.analysis | Analysis/
Cetus.transforms Openacc.transforms | transformation
 | passes

Cetus.exec Openacc.exec | Main driver for
 | parser and
 | various passes

Cetus.codegen Openacc.codegen | Custom
 | codegen pass

OpenARC Main Driver

- openacc.exec.ACC2GPUDriver.java

```
public void run(Strings[] args)
{
    parseCommand(args);
    parseGpuConfFile();
    parseFiles();
    setPasses();
    runPasses();
    CodeGenPass.run(new acc2gpu(...));
    program.print();
}
```

-gpuConfFile=conf.txt
-verbosity=1

-AccAnalysisOnly
-AccPrivatization=1
-AccReduction=2

foo.c
bar.c

Set and run standard analysis/transformation

Performs analysis/transformations for OpenACC-to-GPU translation

cetus_output/foo.cpp
cetus_output/openarc_kernel.cu/cl

<http://ft.ornl.gov/re>

OpenACC-to-GPU code generator

- openacc.codegen.acc2gpu.java

```
public void start()
{
    read command-line options and set parameters
    AccAnnotationParser()
    KernelCallingProcCloning()
    AccLoopDirectivePreprocessor()
    AccAnalysis()
    AccPrivatization/AccReduction()
    Various analyses to be added here.
    ...
}
```

OpenACC-to-GPU code generator (2)

- openacc.codegen.acc2gpu.java

```
//latter part of public void start()
```

```
...
```

```
KernelsSplitting()
```

```
CompRegionConfAnalysis()
```

```
CollapseTransformation()
```

```
ACC2GPUTranslator()
```

```
renameOutputFiles()
```

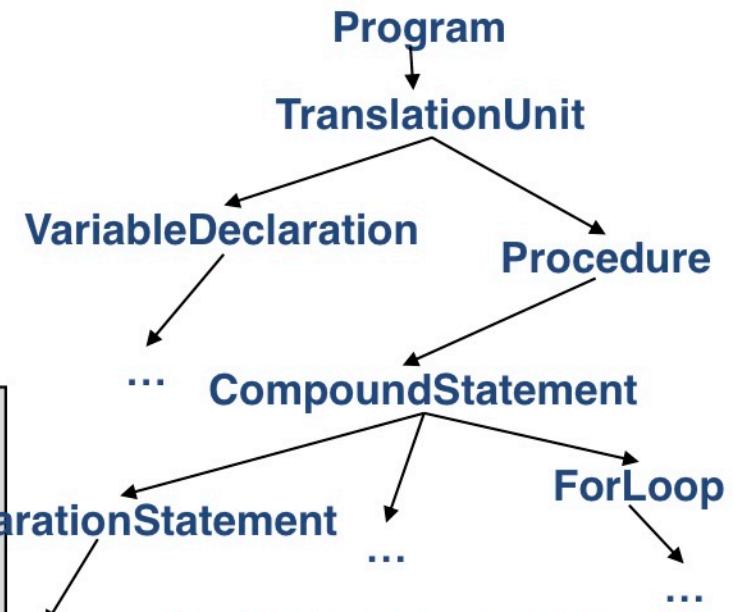
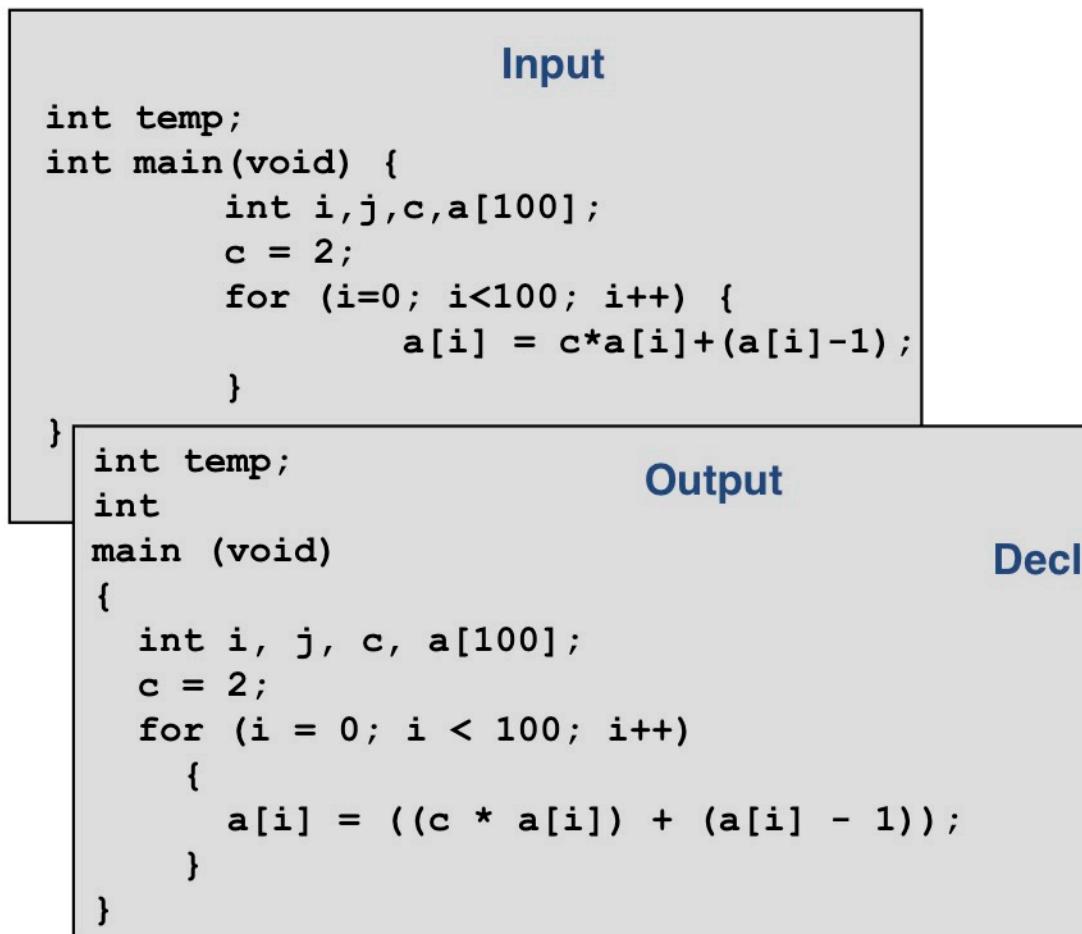
```
}
```

Depending on the target architecture, different derived classes are invoked.

```
GPUInitializer()  
handleDeclareDirectives()  
handleDataRegions()  
convComputeRegions()  
handleUpdateDirectives()  
handleHostDataDirectives()  
handleWaitDirectives()
```

Example of “Built-in” OpenARC Functionality

- OpenARC source-to-source translation – “parse and print”



Simple Tree-Based IR

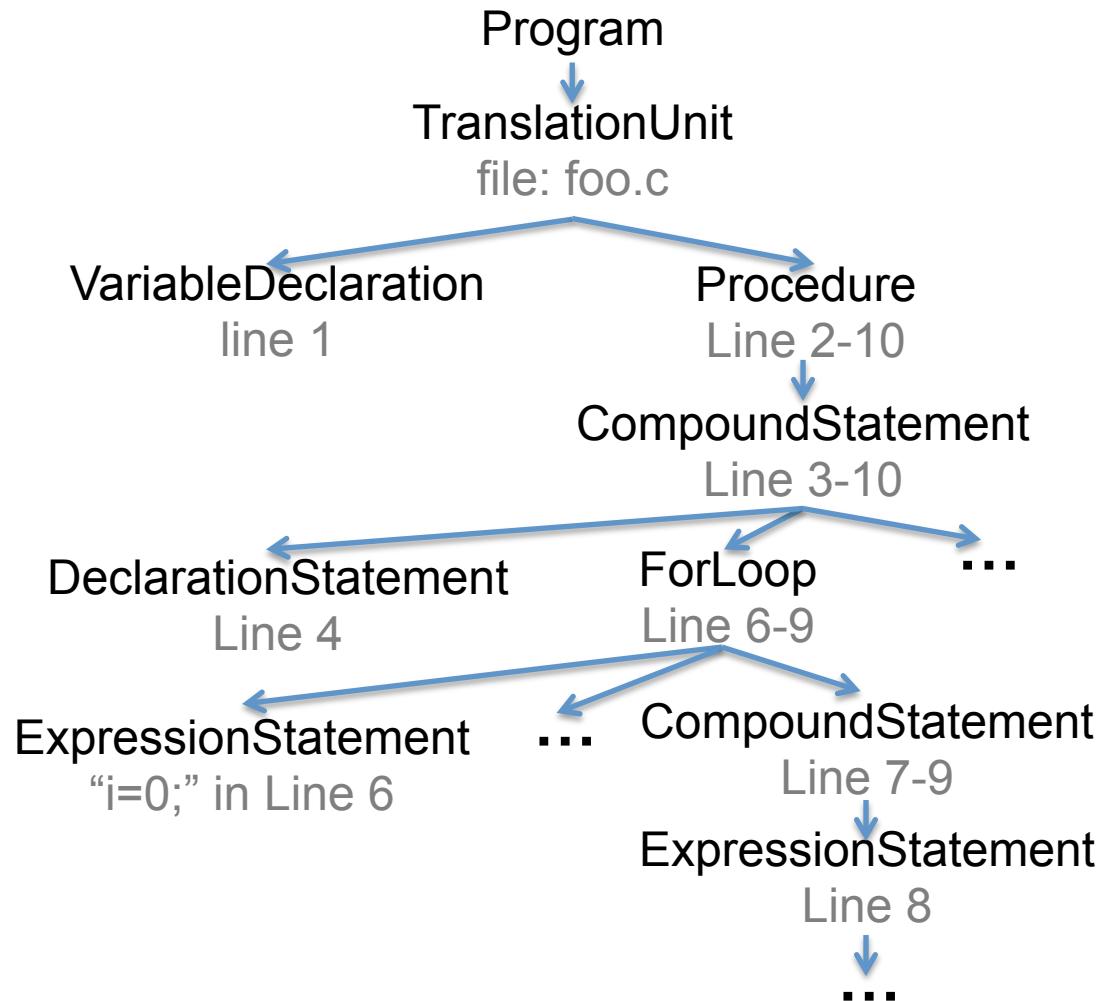
As closely associated with
original program structure as
possible for regeneration of
source code

Session 3: OpenARC Internals (Advanced Topics)

- OpenARC Implementation
- OpenARC Internal Representation (IR)
 - OpenARC IR class hierarchy
 - Working on OpenARC IR
 - Example Transformations
- Built-in Cetus Passes for Program Analysis and Transformation

OpenARC High-Level Intermediate Representation Example

```
0 /* file: foo.c */  
1 int c = 10;  
2 int main(void)  
3 {  
4     int i, a[100], b[100];  
5     #pragma acc parallel loop  
6     for( i=0; i<100; i++ )  
7     {  
8         a[i] = c*b[i];  
9     }  
10    ...  
11 }
```



Major Parts of Class Hierarchy

Base Classes

Program
TranslationUnit
Declaration

Procedure
VariableDeclaration
ClassDeclaration
AnnotationDeclaration
...

Statement

CompoundStatement
ForLoop
ExpressionStatement
AnnotationStatement
...

Expression

BinaryExpression
FunctionCall
...

Sub Classes

Base Classes

Declarator

Sub Classes

VariableDeclarator
ProcedureDeclarator
...

Annotation

CommentAnnotation
CodeAnnotation
PragmaAnnotation

CetusAnnotation
OmpAnnotation
ACCAAnnotation
ASPENAnnotation
...

BinaryOperator

AssignmentOperator
AccessOperator

Specifier

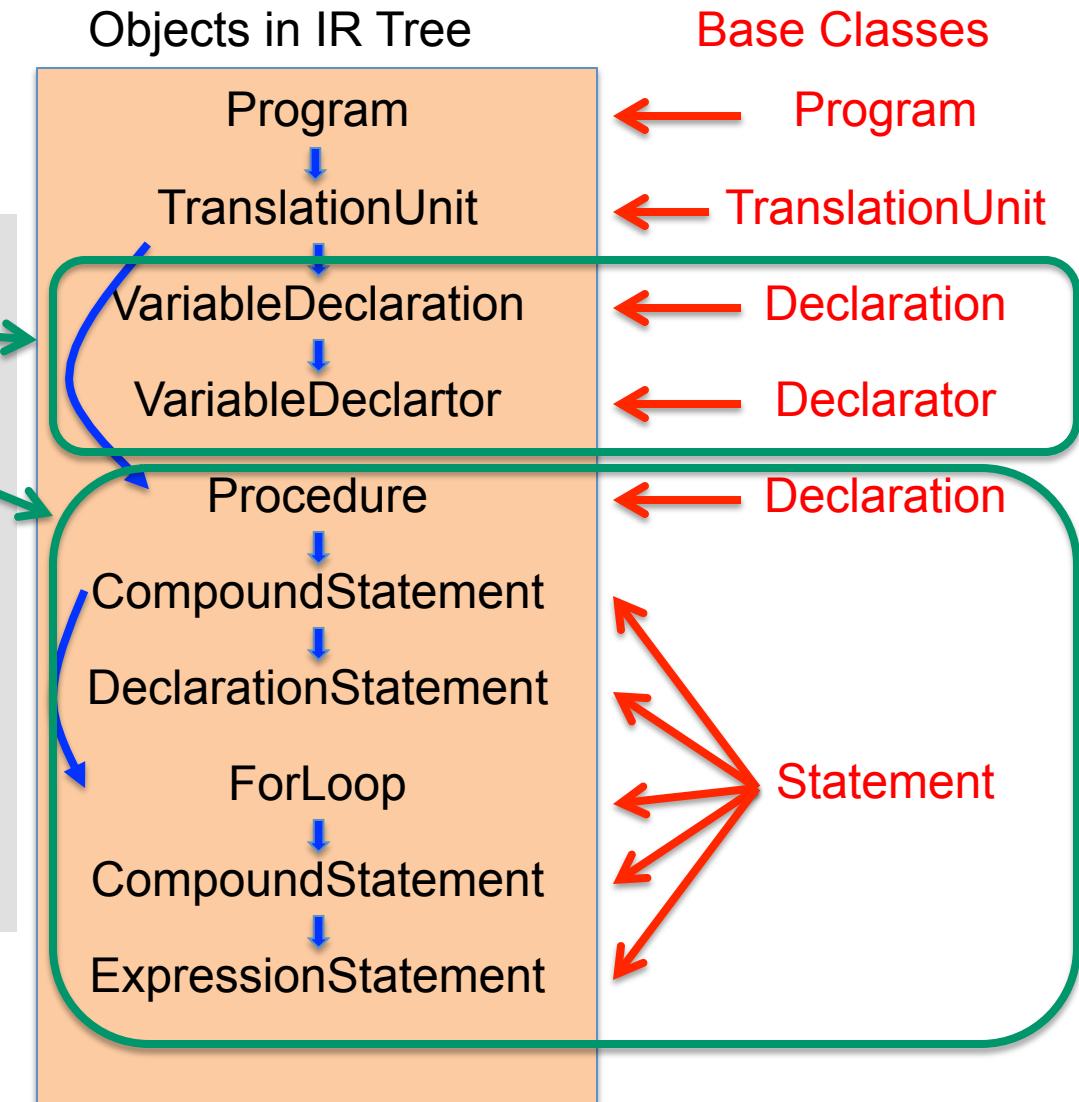
PointerSpecifier
...

Traversable
Not
traversable

OpenARC IR Tree versus Class Hierarchy

→ IR/Syntax Tree
→ Class Hierarchy

```
0 /* file: foo.c */  
1 int c = 10;  
2 int main(void)  
3 {  
4     int i, a[100], b[100];  
5     #pragma acc parallel loop  
6     for( i=0; i<100; i++ )  
7     {  
8         a[i] = c*b[i];  
9     }  
10    ...  
11 }
```

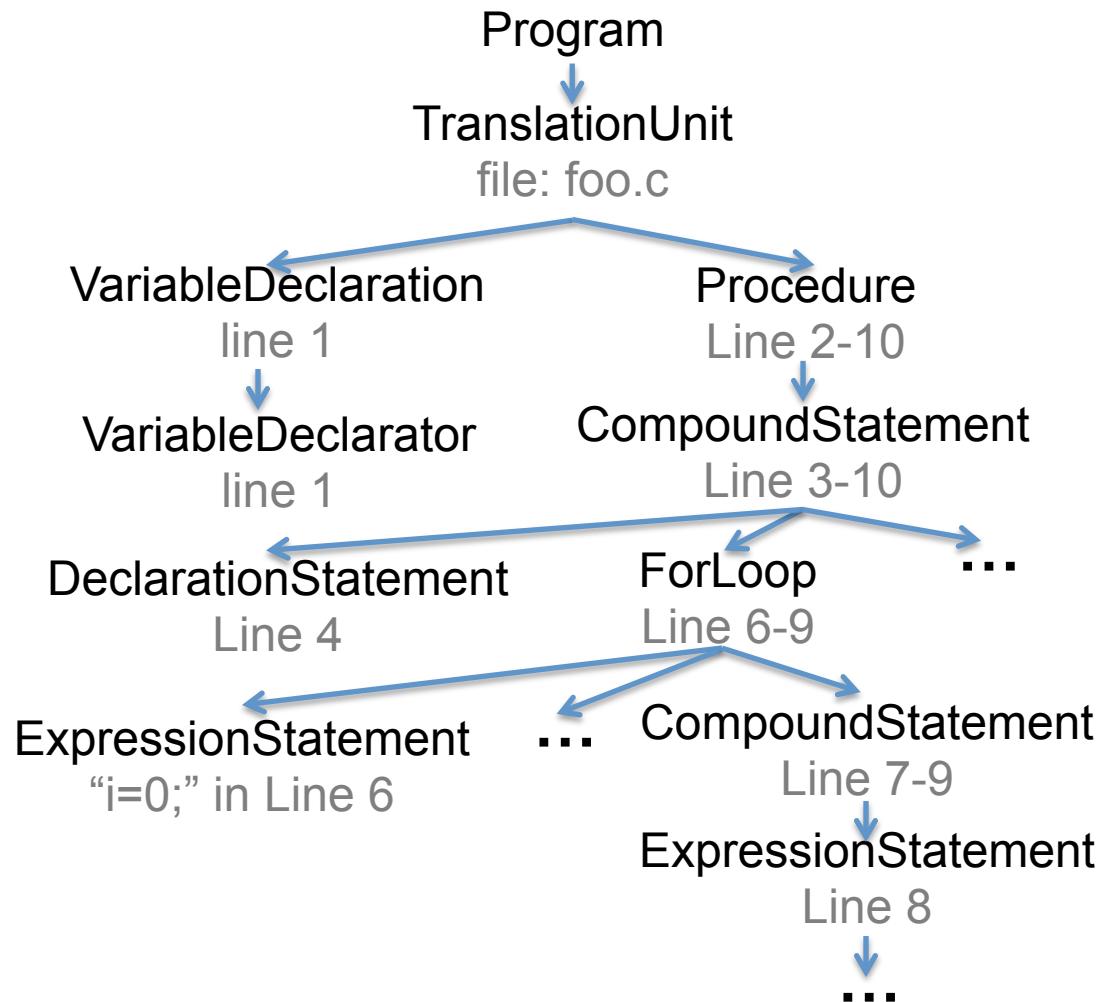


Working on OpenARC IR

- IR Iterators

- DepthFirstIterator
- BreadthFirstIterator
- FlatIterator
- PostOrderIterator

```
0 /* file: foo.c */
1 int c = 10;
2 int main(void)
3 {
4     int i, a[100], b[100];
5     #pragma acc parallel loop
6     for( i=0; i<100; i++ )
7     {
8         a[i] = c*b[i];
9     }
10 ...
11 }
```

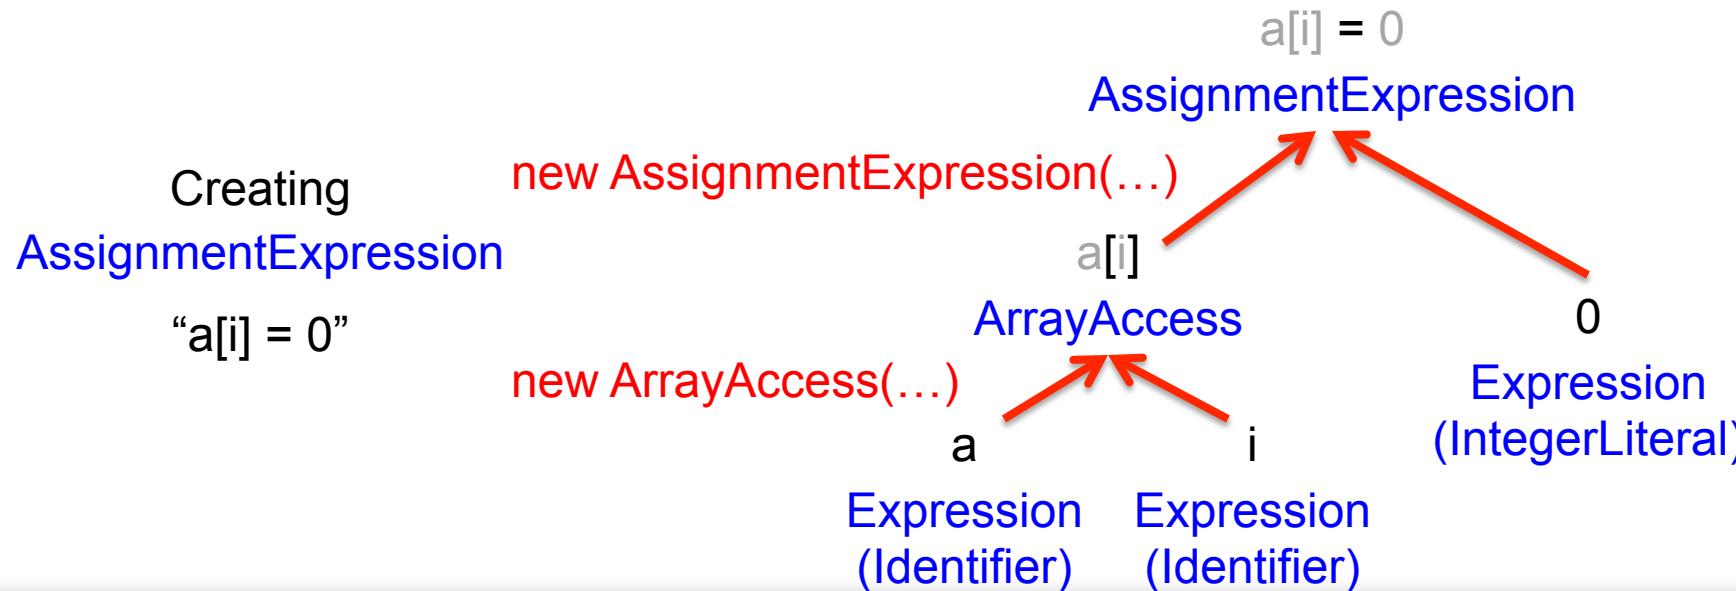


Other Tools assisting IR Traversing

- CFGraph: supports creation of statement-level control flow graphs.
 - Nodes in a CFGraph object contain information gathered from their corresponding OpenARC IR.
- CallGraph: create a static call graph for the program.
- And many others ...

Creating OpenARC IR Objects

- Equivalent to building a tree from leaves
 - Some IRs are created first then populated.
 - E.g., CompoundStatement
 - IR constructors perform the task.
 - Use cloning if creating from an existing IR object.



Modifying OpenARC IR Objects

- Equivalent to inserting/replacing a node in a tree
- IR methods handle most modifications.
- Other modifications are usually search/replace.
 - Create a new expression or statement.
 - Traverse the IR to locate the position for the new node to be inserted or replaced with.
 - Perform insertion or replacement.

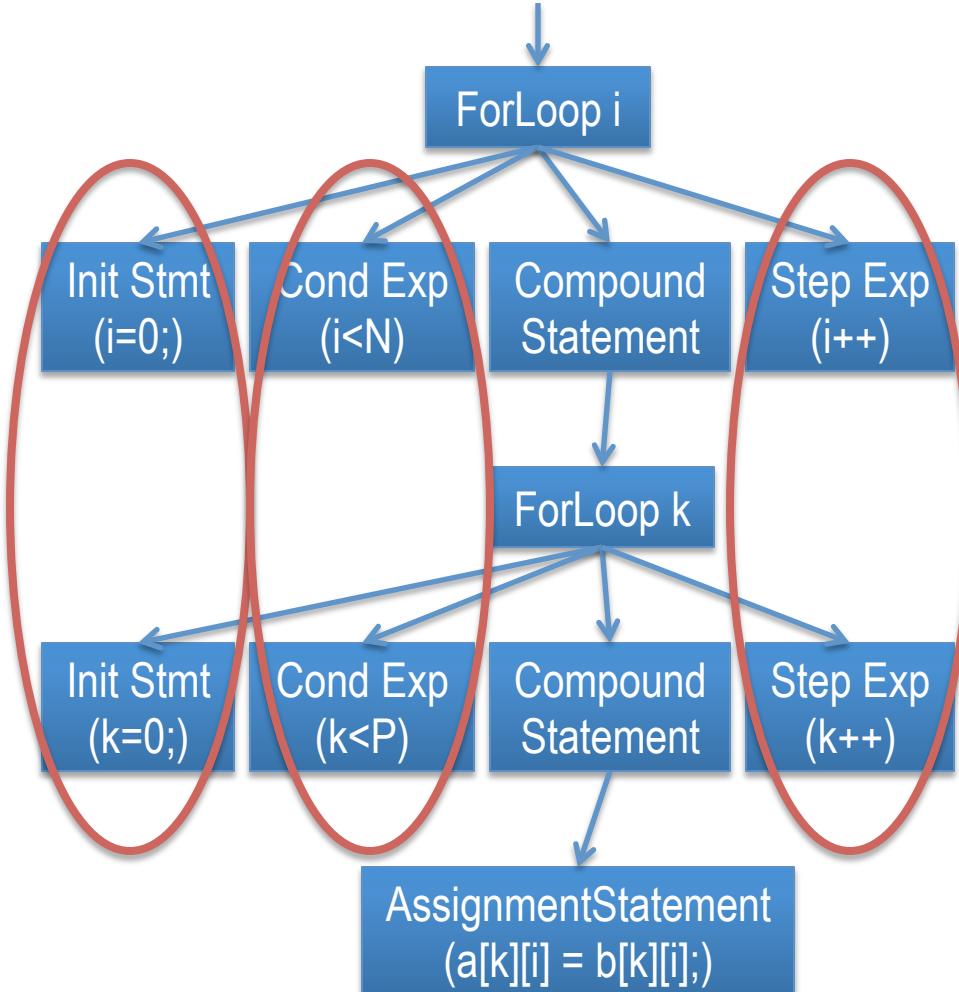
Example Transformation – Loop swap

```
for(i=0; i<N; i++) {  
    for(k=0; k<P; k++) {  
        a[k][i] = b[k][i];  
    }  
}
```



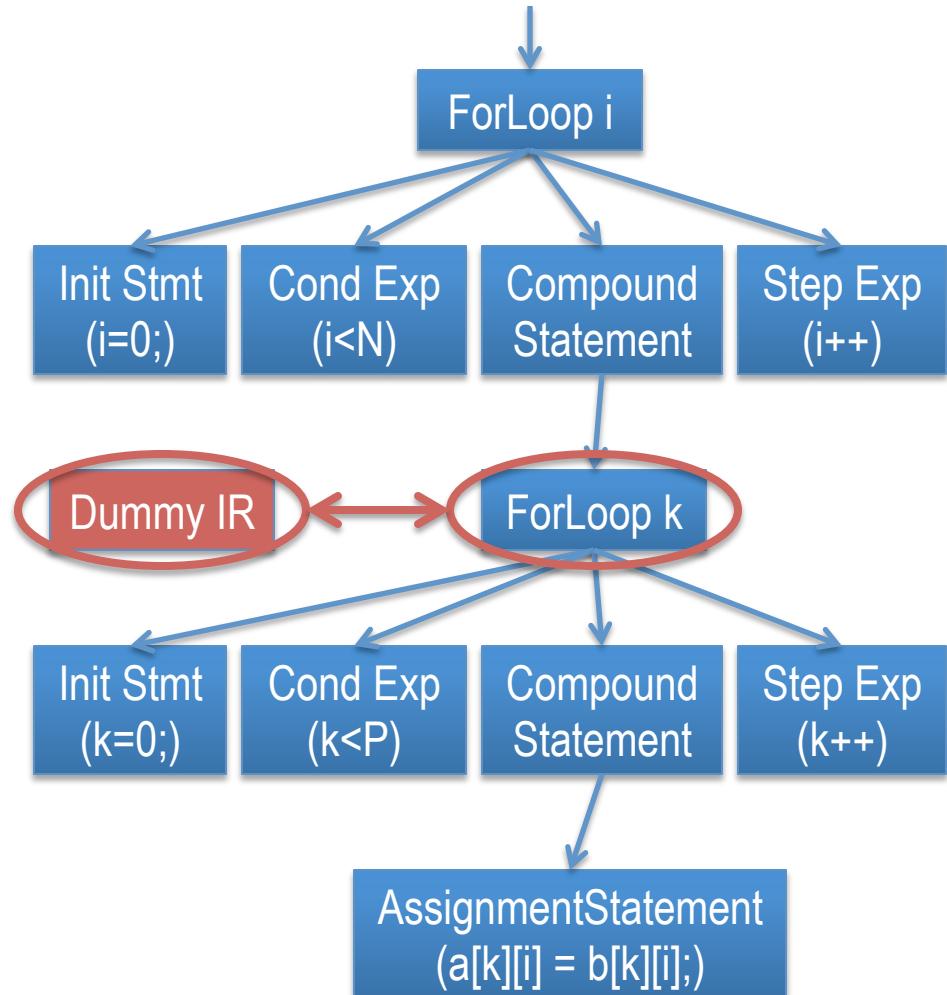
Swap each
child except
loop body

```
for(k=0; k<P; k++) {  
    for(i=0; i<N; i++) {  
        a[k][i] = b[k][i];  
    }  
}
```



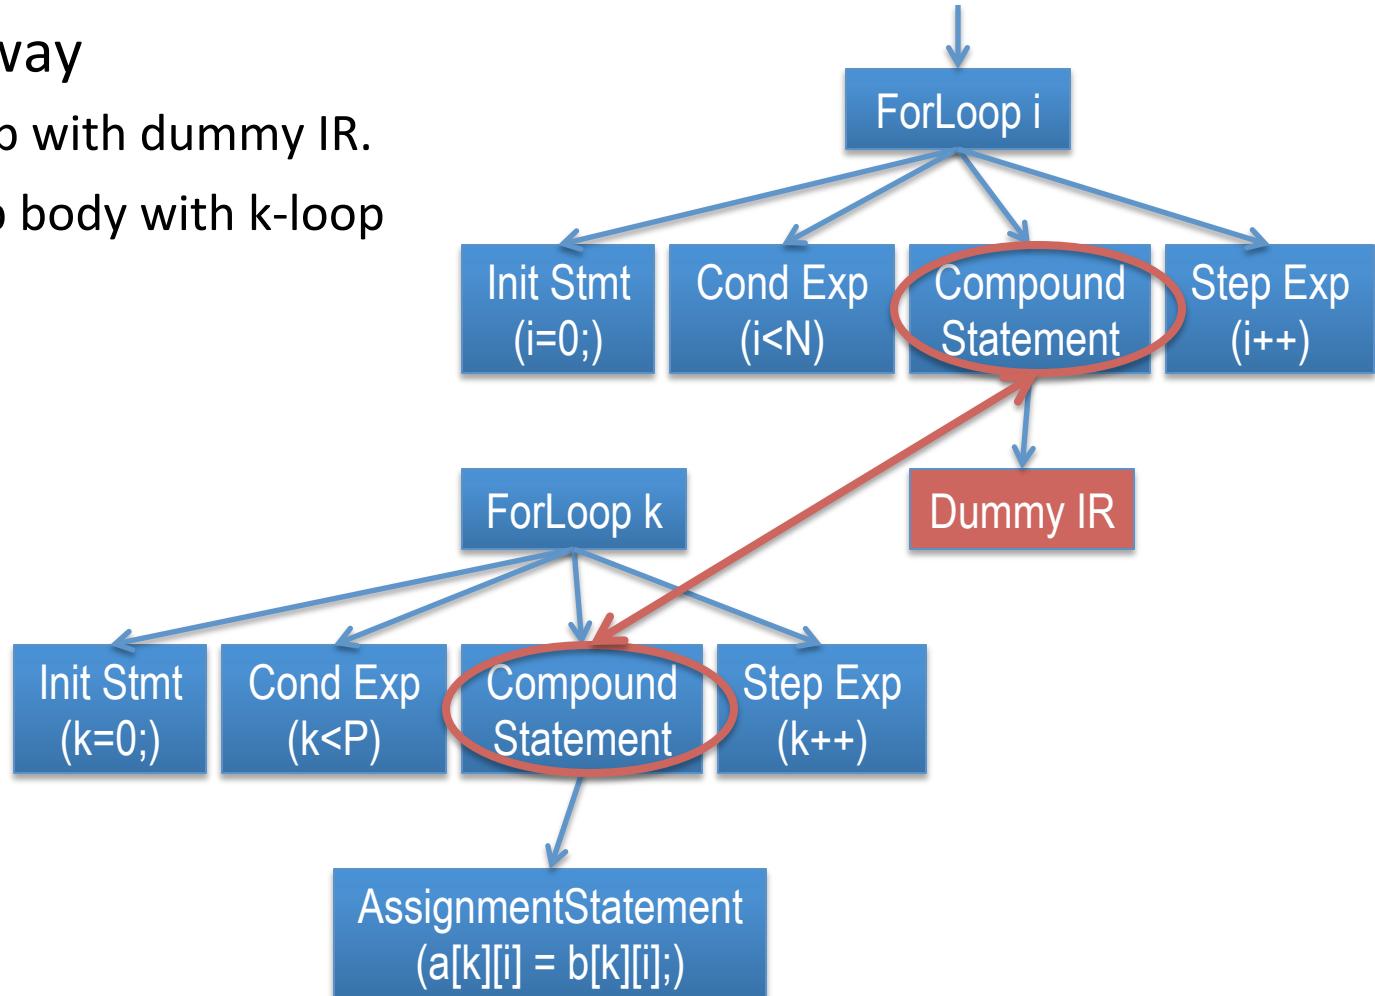
Example Transformation – Loop swap (cont.)

- Alternative way
1) Swap k-loop with dummy IR.



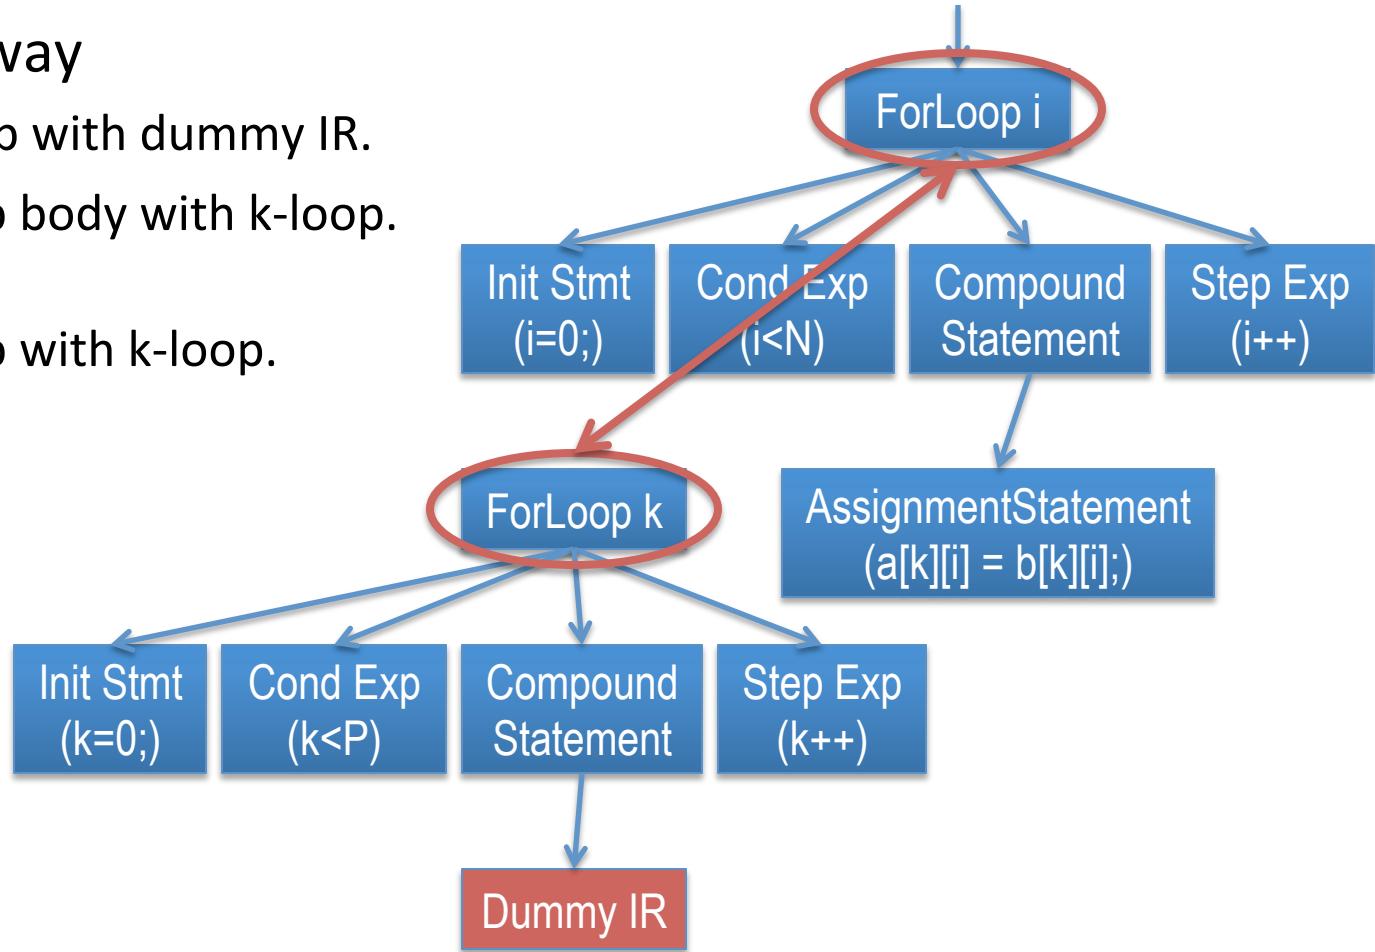
Example Transformation – Loop swap (cont.)

- Alternative way
 - 1) Swap k-loop with dummy IR.
 - 2) Swap i-loop body with k-loop body.



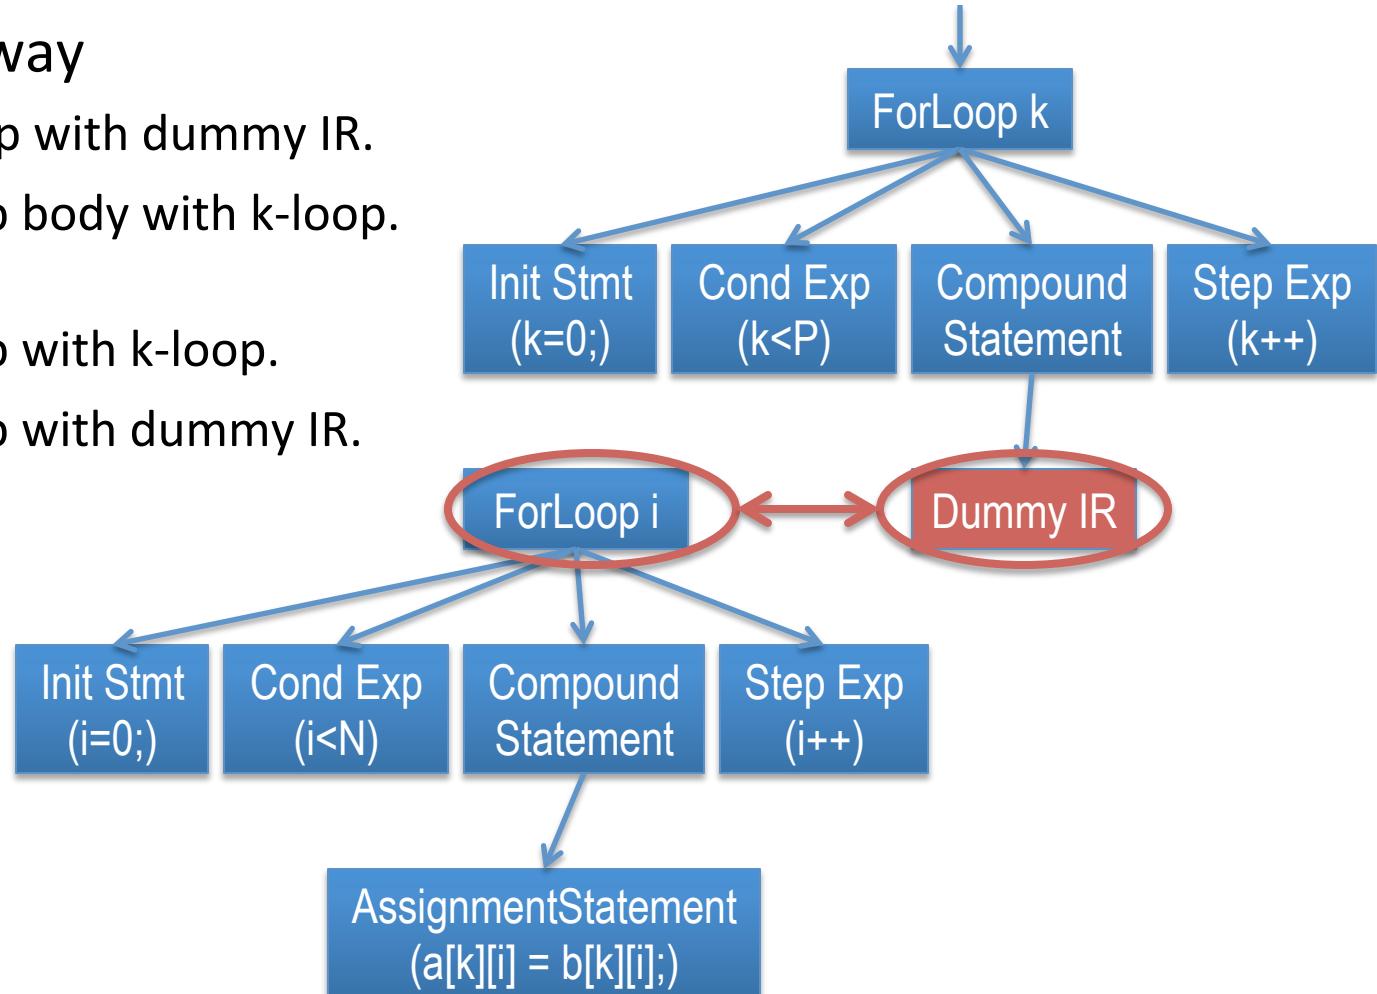
Example Transformation – Loop swap (cont.)

- Alternative way
 - 1) Swap k-loop with dummy IR.
 - 2) Swap i-loop body with k-loop Body.
 - 3) Swap i-loop with k-loop.



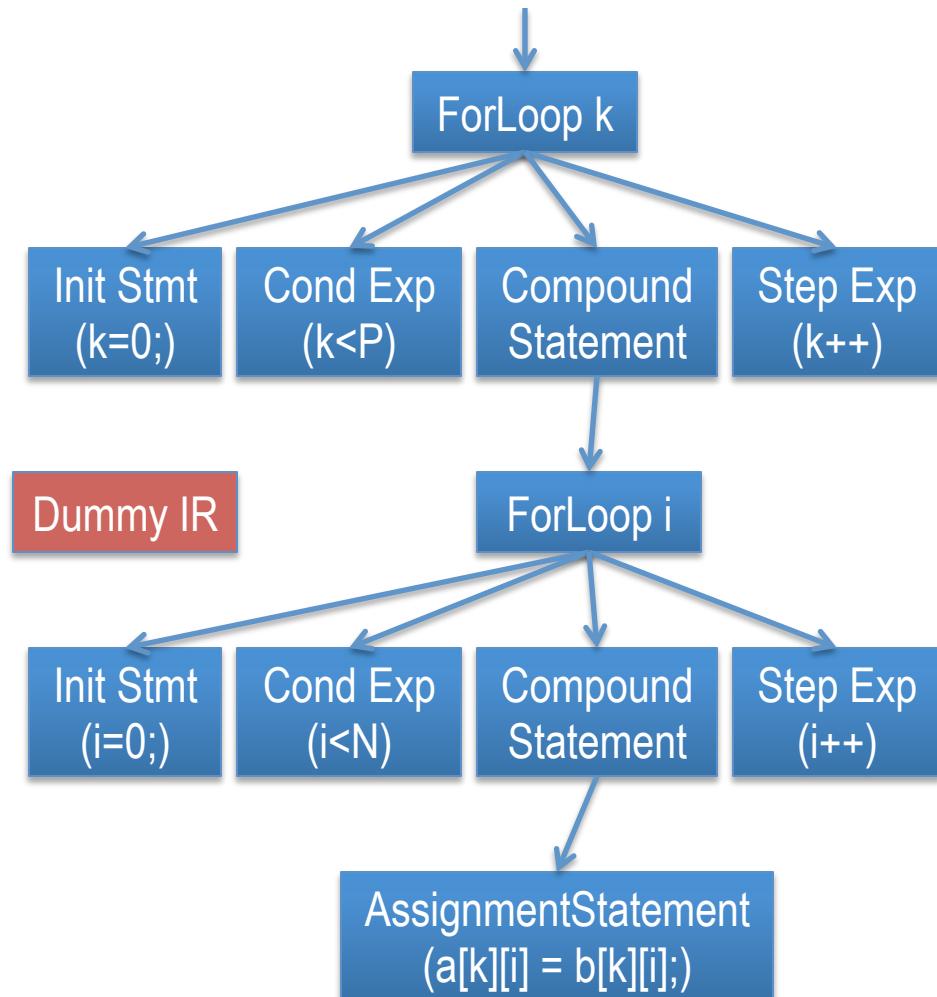
Example Transformation – Loop swap (cont.)

- Alternative way
 - 1) Swap k-loop with dummy IR.
 - 2) Swap i-loop body with k-loop body.
 - 3) Swap i-loop with k-loop.
 - 4) Swap i-loop with dummy IR.



Example Transformation – Loop swap (cont.)

- Alternative way
 - 1) Swap k-loop with dummy IR.
 - 2) Swap i-loop body with k-loop body.
 - 3) Swap i-loop with k-loop.
 - 4) Swap i-loop with dummy IR.



Example Transformation 2: Loop Unrolling

- Loop Unrolling by Factor of 2

- 1) Find loop index, I
- 2) Create expression $i+=2$
- 3) Replace step expression ($i++$) with $i+=2$

```
For (i=0; i<10; i++)  
{  
    fx[i] = fx[i] + x;  
    fy[i] = fy[i] + y;  
    fz[i] = fz[i] + z;  
}
```



```
For (i=0; i<10; i+=2)  
{  
    fx[i] = fx[i] + x;  
    fy[i] = fy[i] + y;  
    fz[i] = fz[i] + z;  
    fx[i+1] = fx[i+1] + x;  
    fy[i+1] = fy[i+1] + y;  
    fz[i+1] = fz[i+1] + z;  
}
```

- 1) Copy the statements in the loop body.
- 2) Find loop index, i
- 3) Search & replace expression i with $i+1$
- 4) Insert the modified statements

Creating/Inserting Annotations

- Annotations are used in OpenARC for
 - Information exchange between passes - PragmaAnnotation
 - E.g., set of modified/used variables
 - Information for backend compilers – PragmaAnnotation
 - E.g., OpenMP pragmas, OpenACC pragmas
 - Information for code readers – CommentAnnotation
 - Inline arbitrary codes by the compiler – CodeAnnotation
- Annotation can be either
 - Stand-alone annotation or
 - Attached annotation: associated with a specific statement or a declaration

Creating/Inserting Annotations (cont.)

- Steps
 - Create an annotation with an appropriate type
 - CommentAnnotation, PragmaAnnotation, ...
 - Locate the position for the annotation
 - Attached annotation: the associated statement/declaration
 - Stand-alone annotation: the reference statement/declaration
 - Insert the annotation
 - Attached annotation: insert directly to the statement/declaration
 - Use Annotatable interface, which supports annotate(), getAnnotation(), ...
 - All Declaration/Statement IRs implement the Annotatable interface.
 - Stand-alone annotation: encapsulate in a dummy statement/declaration and insert around the reference.
 - Use AnnotationStatement in a CompoundStatement or
 - AnnotationDeclaration in a TranslationUnit

Creating Custom PragmaAnnotations

- Option 1: use existing PragmaAnnotation
 - Annotation class extends HashMap, and thus it can keep arbitrary (key, value) pairs.
 - Useful for internal use among compiler passes
- Option2: derive a custom Annotation class derived from PragmaAnnotation.
 - May need to provide additional annotation parser to automatically attach the annotations to other IRs.
 - Default C parser creates all annotations as stand-alone type.
 - Additional parsers (OmpParser, ACCParser) are used to reason their semantics and attach them to related IRs accordingly.

Session 3: OpenARC Internals (Advanced Topics)

- OpenARC Implementation
- OpenARC Internal Representation (IR)
 - OpenARC IR class hierarchy
 - Working on OpenARC IR
 - Example Transformations
- Built-in Cetus Passes for Program Analysis and Transformation

Analysis Passes

- Data dependence analysis
 - Banerjee-Wolfe Test
 - Range Test
- Pointer alias analysis
- Symbolic range analysis
 - Generates a map from variables to their valid symbolic bounds at each program point.
- Array privatization
 - Compute privatizable scalars and arrays
- Reduction recognition
 - Detects reduction operations

Analysis Passes (cont.)

- Symbolic expression manipulators
 - Normalizes and simplifies expressions.
 - Examples
 - $1+2*a+4-a \rightarrow 5+a$ (folding) $(a^2)/(8*c) \rightarrow a/(4*c)$ (division)
- Call Graph and CFG generators
 - CFG provided either at basic-block level or at statement level
- Basic use/def set computation for both scalars and array sections
- And many others...

Transformation Passes

- Induction variable substitution pass
 - Identifies and substitutes induction variables.
- Loop parallelizer
 - Depends on induction variable substitution, reduction recognition, and array privatization.
 - Performs loop dependence analysis.
 - Generates “parallel loop” annotations.
- Program normalizers
 - Single call, single declarator, single return normalizers.
- Loop outliner
 - Extract loops out into separate subroutines.
- And many others...

Thank You!

For more information, please refer to
the OpenARC website
<http://ft.ornl.gov/research/openarc>

